

Analysis of Cross-Sectional Data

Kevin Sheppard

Course Structure

- Course presented through three channels:
 1. Pre-recorded content with a focus on technical aspects of the course
 - Designed to be viewed in sequence
 - Each module should be short
 - Approximately 2 hours of content per week
 2. In-person lectures with a focus on applied aspects of the course
 - Expected that pre-recorded content has been viewed before the lecture
 3. Notes that accompany the lecture content
 - Read before or after the lecture or when necessary for additional background
- Slides are primary – material presented during lecturers, either pre-recorded or live is examinable
- Notes are secondary and provide more background for the slides
- Slides are derived from notes so there is a strong correspondence

Monitoring Your Progress

- Self assessment
 - Review questions in pre-recorded content
 - Multiple choice questions on Canvas made available each week
 - Answers available immediately
 - Long-form problem distributed each week
 - Answers presented in a subsequent class
- Marked Assessment
 - Empirical projects applying the material in the lectures
 - Both individual and group
 - Each empirical assignment will have a written and code component

Analysis of Cross-Sectional Data

Introduction to Regression Models

- Notation
- Factor Models
- Data
- Variable Transformations

Linear Regression

Scalar Notation

$$Y_i = \beta_1 X_{1,i} + \beta_2 X_{2,i} + \dots + \beta_k X_{k,i} + \epsilon_i$$

- Y_i : Regressand, Dependent Variable, LHS Variable
- $X_{j,i}$: Regressor, also Independent Variable, RHS Variable, Explanatory Variable
- ϵ_i : Innovation, also Shock, Error or Disturbance
- n observations, indexed $i = 1, 2, \dots, n$
- k regressors, indexed $j = 1, 2, \dots, k$

Linear Regression

Matrix Notation

Common to use convenient matrix notation

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

- \mathbf{y} is n by 1
- \mathbf{X} is n by k
- $\boldsymbol{\beta}$ is k by 1
- $\boldsymbol{\epsilon}$ is n by 1

Factor Models

- Factor models are widely used in finance
 - Capital Asset Pricing Model (CAPM)
 - Arbitrage Pricing (APT)
 - Risk Exposure
- Basic specification $R_i = \mathbf{f}_i \boldsymbol{\beta} + \epsilon_i$
 - R_i : Return on dependent asset, often excess (R_i^e)
 - \mathbf{f}_i : $1 \times k$ vector of factor innovations
 - ϵ_i innovation, $\text{corr}(\epsilon_i, F_{j,i}) = 0, j = 1, 2, \dots, k$
- Special Case: CAPM

$$\begin{aligned}R_i - R_i^f &= \beta(R_i^m - R_i^f) + \epsilon_i \\R_i^e &= \beta R_i^{me} + \epsilon_i\end{aligned}$$

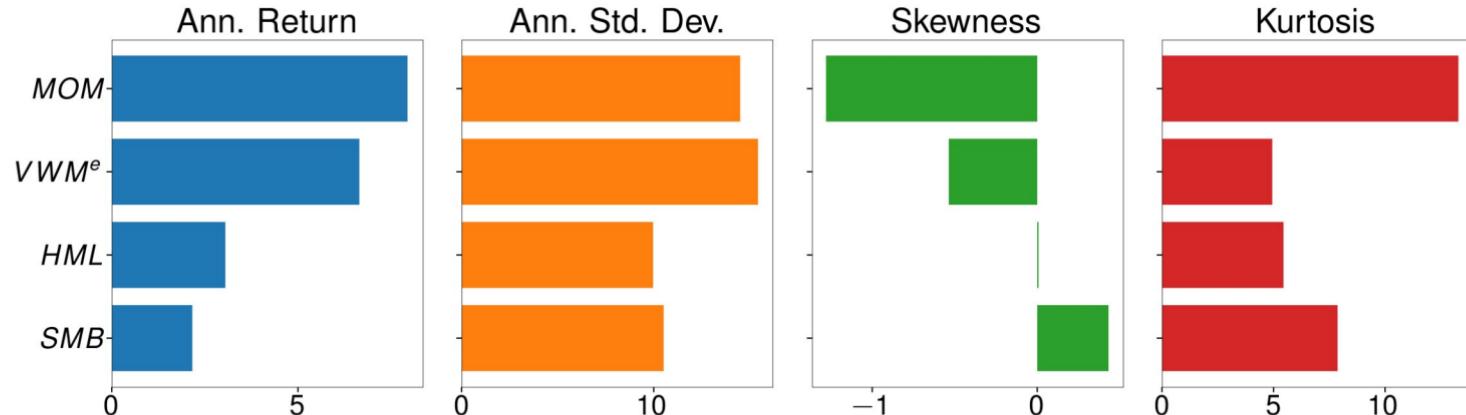
Data

- Data from the Fama-French 3 factors + Momentum
 - VWM^e - Excess return on Value-Weighted-Market
 - SMB - Return on the size portfolio
 - HML - Return on the value portfolio
 - MOM - Return on the momentum portfolio
- Size-Value sorted portfolio return data
 - Size
 - S: Small
 - B: Big
 - Value
 - H: High BE/ME
 - M: Middle BE/ME
 - L: Low BE/ME
- 49 Industry Portfolios
- All returns excess except SMB, HML, MOM

Fama-French Factors

Summary Statistics

```
In [3]: summ_plot(factors)
```



Fama-French Factors

Correlation Structure

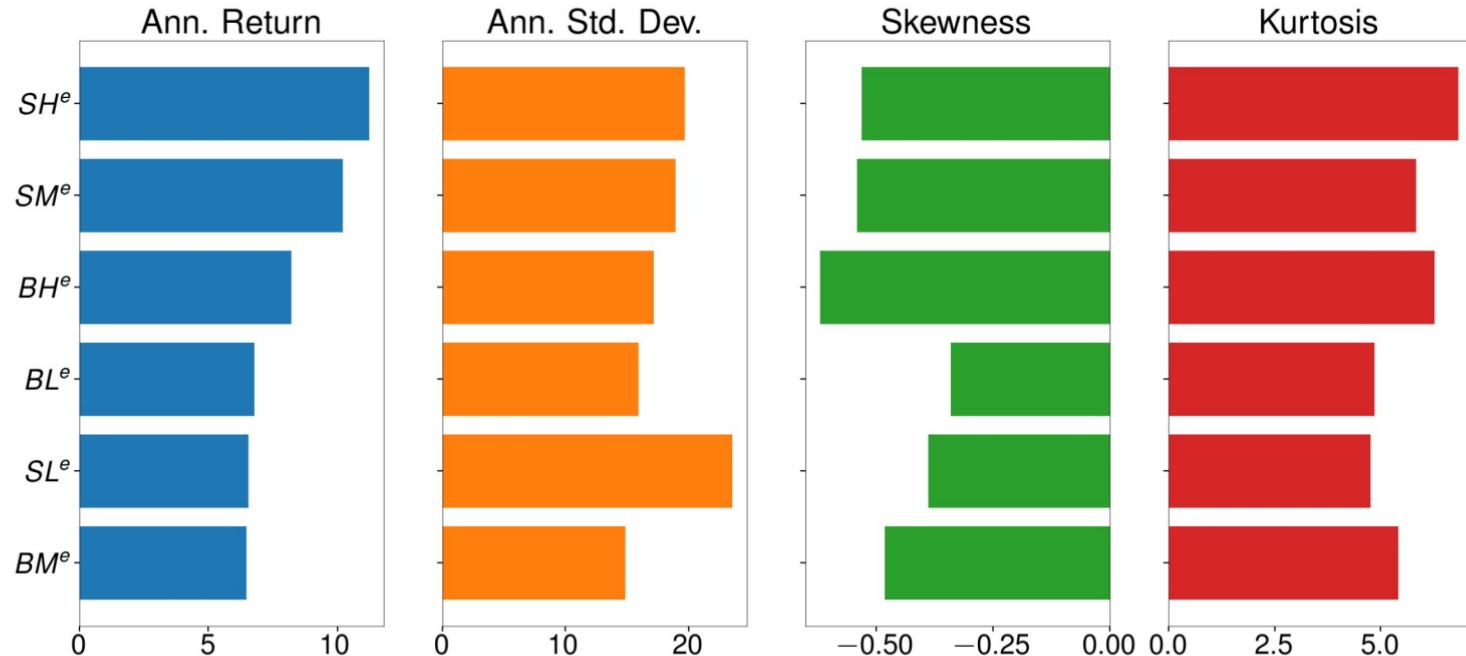
```
In [4]: factors.corr()
```

```
Out[4]:
```

	<i>VWM^e</i>	<i>SMB</i>	<i>HML</i>	<i>MOM</i>
<i>VWM^e</i>	1.000000	0.300958	-0.226222	-0.149518
<i>SMB</i>	0.300958	1.000000	-0.174962	-0.024014
<i>HML</i>	-0.226222	-0.174962	1.000000	-0.195242
<i>MOM</i>	-0.149518	-0.024014	-0.195242	1.000000

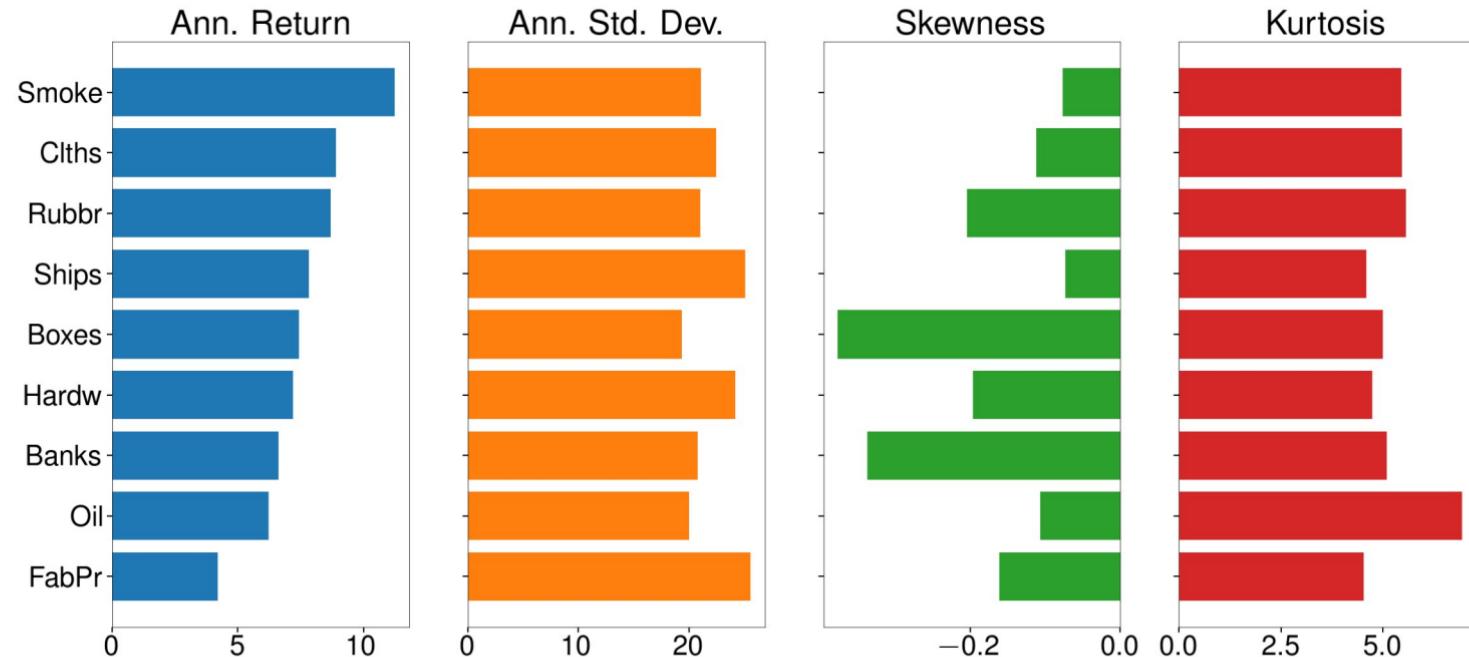
Size and Value components

```
In [6]: summ_plot(components)
```



Industry Portfolios

```
In [7]: summ_plot(subset)
```



Variable Transformations

- Dummy variables
 - 0-1 variables based on an indicator function

$$I_{[X_{i,j} > 0]}$$

- Asymmetries at 0
 - Monthly Effects

```
In [9]: monthly_dummies.head(8)
```

Out[9]:

Variable Transformation: Interactions

- Interactions dramatically expand the functional forms that can be specified
 - Powers and Cross-products: $X_{i,j}^2, X_{i,j}X_{i,m}$
 - Dummy Interactions to Produce Asymmetries: $X_{i,j} \times I_{[X_{i,j} < 0]}$

```
In [11]: interactions.tail(10)
```

Out[11]:

	Market Negative	Negative Return	Squared Returns
2019-11-30	0	0.00	14.9769
2019-12-31	0	0.00	7.6729
2020-01-31	1	-0.11	0.0121
2020-02-29	1	-8.13	66.0969
2020-03-31	1	-13.38	179.0244
2020-04-30	0	0.00	186.3225
2020-05-31	0	0.00	31.1364
2020-06-30	0	0.00	6.0516
2020-07-31	0	0.00	33.2929
2020-08-31	0	0.00	58.0644

Analysis of Cross-Sectional Data

Parameter Estimation and Model Fit

- Parameter Estimation
- Models with Interactions
- Other estimated quantities
- Regression Coefficient in Factor Models

Parameter Estimation

Least Squares

$$\operatorname{argmin}_{\beta} \sum_{i=1}^n (Y_i - \mathbf{x}_i \boldsymbol{\beta})^2$$

```
In [13]: ls = smf.ols("BHe ~ 1 + VWMe + SMB + HML + MOM", data).fit(cov_type="HC0")
summary(ls)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.0859	0.043	-1.991	0.046	-0.170	-0.001
VWMe	1.0798	0.012	93.514	0.000	1.057	1.102
SMB	0.0019	0.017	0.110	0.912	-0.032	0.036
HML	0.7643	0.021	36.380	0.000	0.723	0.805
MOM	-0.0354	0.013	-2.631	0.009	-0.062	-0.009

Parameter Estimation

Least Absolute Deviations

$$\operatorname{argmin}_{\beta} \sum_{i=1}^n |Y_i - \mathbf{x}_i \beta|$$

```
In [14]: lad = smf.quantreg("BHe ~ 1 + VWMe + SMB + HML + MOM", data).fit(q=0.5)
summary(lad)
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-0.0306	0.044	-0.696	0.487	-0.117	0.056
VWMe	1.0716	0.010	103.257	0.000	1.051	1.092
SMB	0.0161	0.015	1.090	0.276	-0.013	0.045
HML	0.7503	0.016	47.702	0.000	0.719	0.781
MOM	-0.0272	0.011	-2.581	0.010	-0.048	-0.007

Estimating Models with Interactions

Added an asymmetry and a square of VWM to the 4-factor model

$$Util_i = \beta_1 + \beta_2 VWM_i^e + \beta_3 (VWM_i^e)^2 + \beta_4 VWM_i^e I_{[VWM_i^e < 0]} + \beta_5 SMB_i + \beta_6 HML_i + \beta_7 MOM_i + \epsilon_i$$

```
In [15]: model = f"Util ~ 1 + VWMe + I(VWMe**2) + I(VWMe * (VWMe < 0)) + SMB + HML + MO  
M"  
ls_interact = smf.ols(model, data).fit(cov_type="HC0")  
summary(ls_interact)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	0.2857	0.225	1.268	0.205	-0.156	0.727
VWMe	0.4594	0.089	5.154	0.000	0.285	0.634
I(VWMe ** 2)	0.0159	0.007	2.240	0.025	0.002	0.030
I(VWMe * (VWMe < 0))	0.3524	0.188	1.870	0.061	-0.017	0.722
SMB	-0.1972	0.048	-4.087	0.000	-0.292	-0.103
HML	0.3470	0.060	5.810	0.000	0.230	0.464
MOM	0.0611	0.039	1.578	0.114	-0.015	0.137

Expected and Fitted Values

- Fitted values:

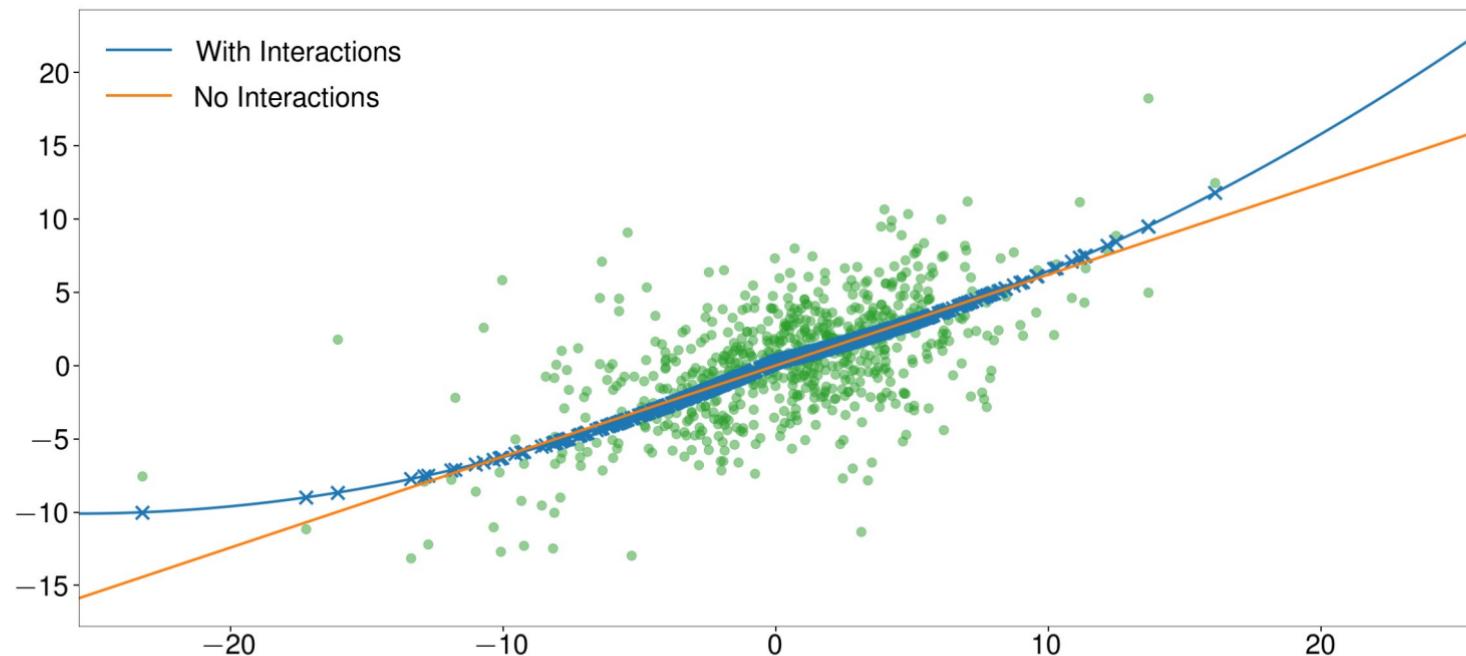
$$\hat{Y}_i = \mathbf{x}_i \hat{\boldsymbol{\beta}}$$

- Expected values:

$$E[Y|X = \mathbf{x}] = \mathbf{x} \hat{\boldsymbol{\beta}}$$

Expected and Fitted Values

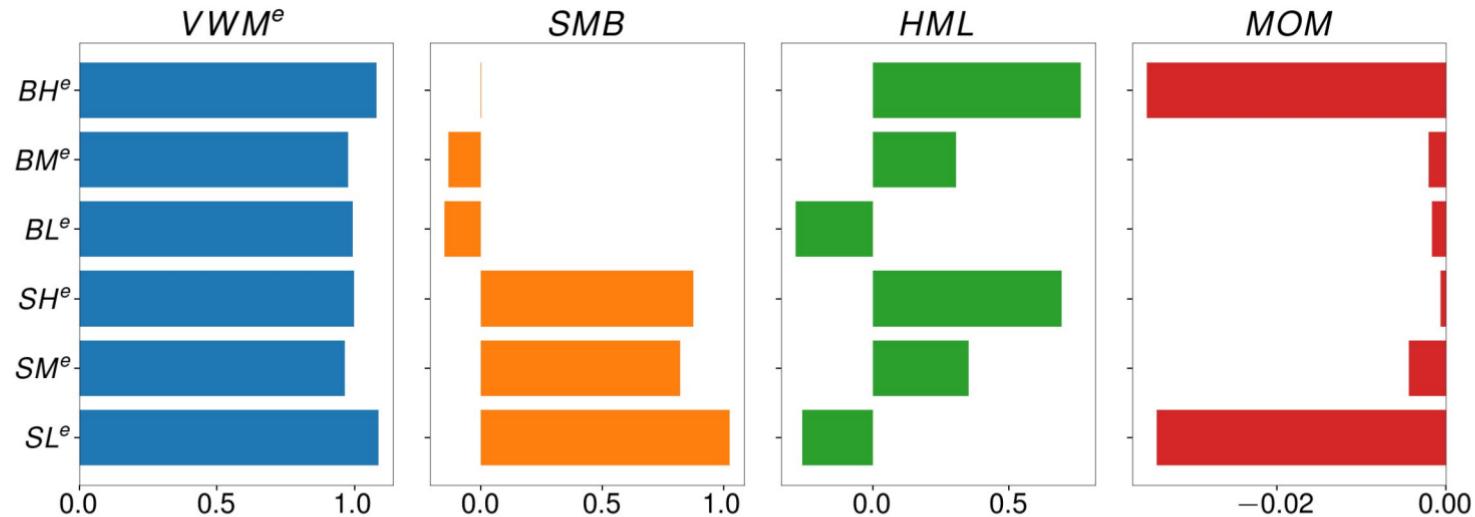
```
In [18]: plot_market_interactions()
```



Typical Regression Coefficients

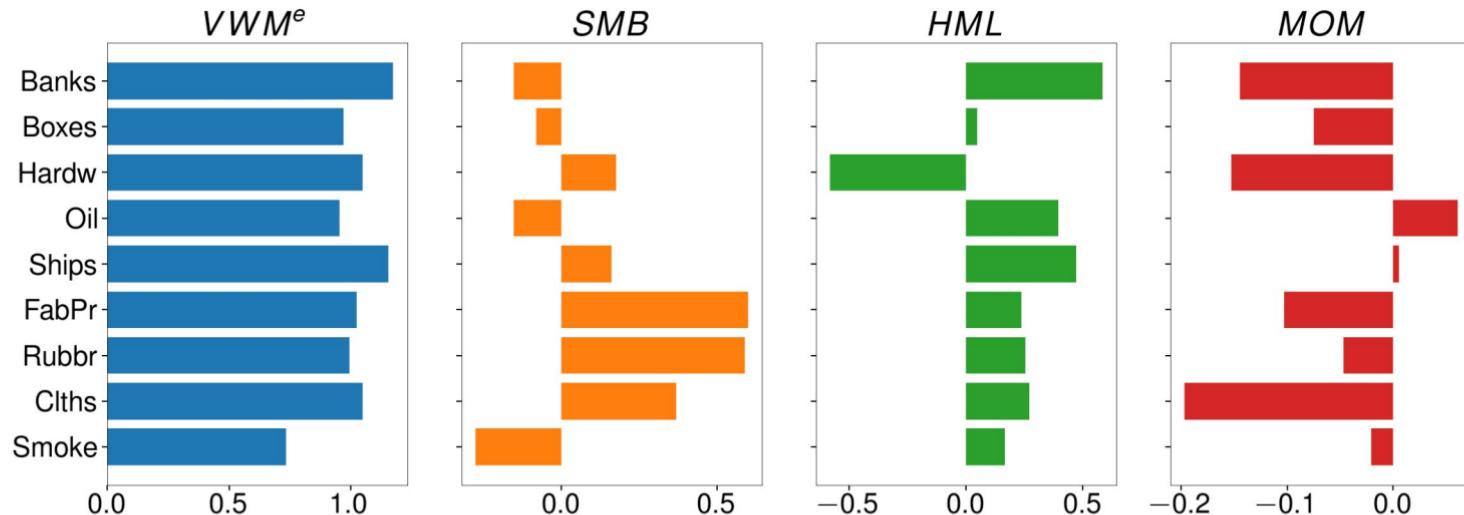
Factor Components

```
In [20]: beta_plot(betas, titles)
```



Typical Regression Coefficients Industry Portfolios

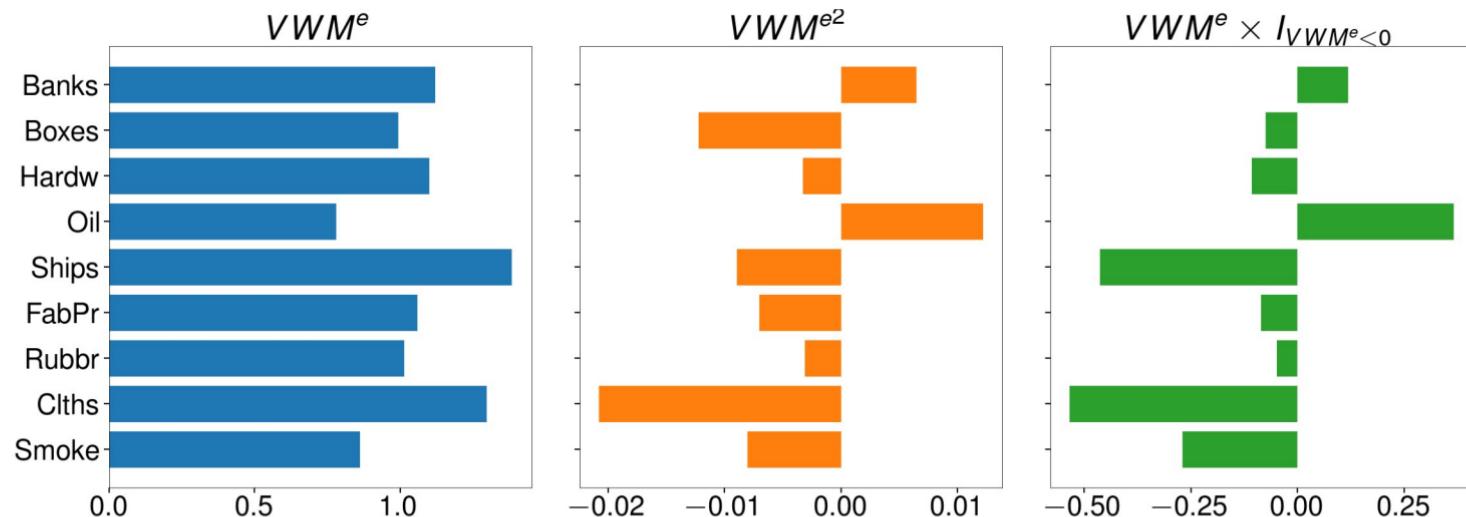
```
In [22]: beta_plot(betas, titles)
```



Evidence of Non-linear returns

- Add square and asymmetry to 4-factor model

```
In [24]: beta_plot(betas, titles)
```



Measuring fit

- Coefficient of Determination

$$R^2 = 1 - \frac{SSE}{TSS} = \frac{RSS}{TSS}$$

- Based on a complete decomposition $TSS = SSE + RSS$
- Total Sum of Squares

$$TSS = \sum_{i=1}^n (Y_i - \bar{Y})^2$$

- Sum of Squared Errors

$$SSE = \sum_{i=1}^n \hat{\epsilon}_i^2$$

- Regression Sum of Squares

$$RSS = \sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2 = \sum_{i=1}^n (\mathbf{x}_i \hat{\beta} - \bar{Y})^2$$

Measuring Fit

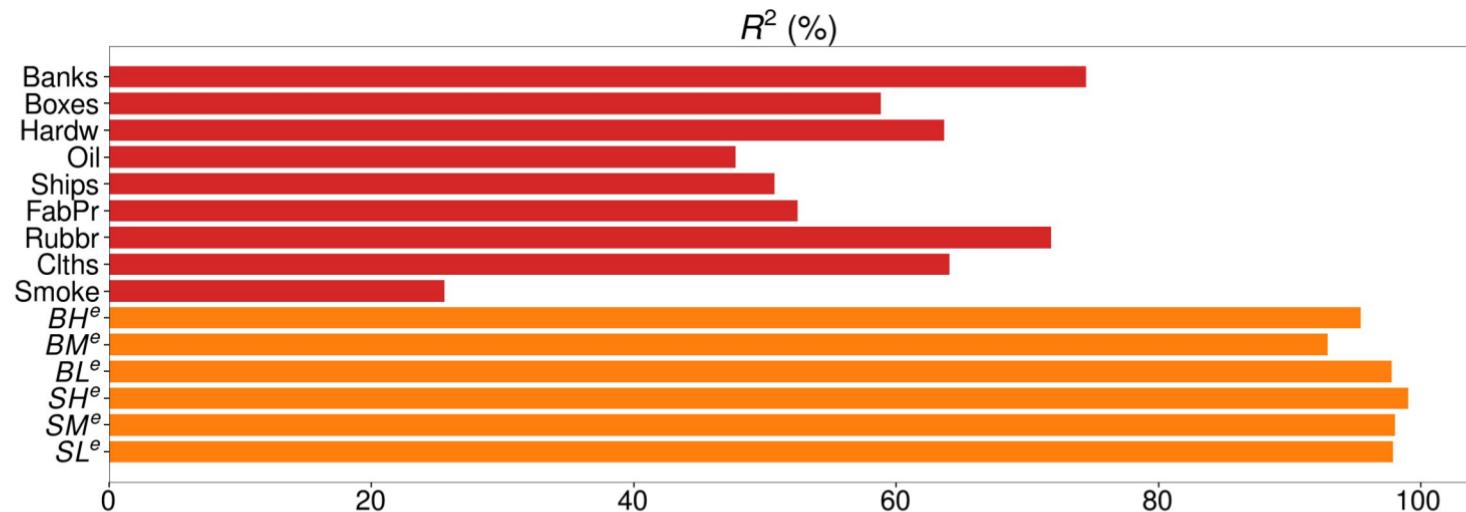
```
In [25]: ls = smf.ols("BHe ~ 1 + VWMe + SMB + HML + MOM", data).fit(cov_type="HC0")
summary(ls, [0])
```

Dep. Variable:	BHe	R-squared:	0.954
Model:	OLS	Adj. R-squared:	0.954

Measuring Fit

Component and Industry Fits

```
In [27]: r2_plot()
```



Measuring Fit

Shifting variables

$$BH_i^e + 99 = \beta_1 + \beta_2 VWM_i^e + \beta_3 SMB_i + \beta_4 HML_i + \beta_5 MOM_i + \epsilon_i$$

```
In [28]: shifted_mod = "I(BHe + 99) ~ 1 + VWMe + SMB + HML + MOM"  
ls_shift = smf.ols(shifted_mod, data).fit(cov_type="HC0")  
summary(ls_shift, [0])
```

Dep. Variable: I(BHe + 99) R-squared: 0.954
Model: OLS Adj. R-squared: 0.954

```
In [29]: summary(ls, [0])
```

Dep. Variable: BHe R-squared: 0.954
Model: OLS Adj. R-squared: 0.954

Measuring Fit

Rescaling variables

$$\pi BH_i^e = \beta_1 + \beta_2 VWM_i^e + \beta_3 SMB_i + \beta_4 HML_i + \beta_5 MOM_i + \epsilon_i$$

```
In [30]: rescaled_mod ="I(np.pi * BHe) ~ 1 + VWMe + SMB + HML + MOM"  
ls_scale = smf.ols(rescaled_mod, data).fit(cov_type="HC0")  
summary(ls_scale, [0])
```

Dep. Variable:	I(np.pi * BHe)	R-squared:	0.954
Model:	OLS	Adj. R-squared:	0.954

```
In [31]: summary(ls, [0])
```

Dep. Variable:	BHe	R-squared:	0.954
Model:	OLS	Adj. R-squared:	0.954

Measuring Fit

Changing the LHS Variable

$$(BH_i^e - VW M_i^e - HML_i) = \beta_1 + \beta_2 VWM_i^e + \beta_3 SMB_i + \beta_4 HML_i + \beta_5 MOM_i + \epsilon_i$$

```
In [32]: model = "I(BHe - VWMe - HML) ~ 1 + VWMe + SMB + HML + MOM"  
ls_lhs = smf.ols(model, data).fit(cov_type="HC0")  
summary(ls_lhs, [0])
```

Dep. Variable: I(BHe - VWMe - HML) R-squared: 0.382
Model: OLS Adj. R-squared: 0.378

```
In [33]: summary(ls, [0])
```

Dep. Variable: BHe R-squared: 0.954
Model: OLS Adj. R-squared: 0.954

Measuring fit

Caveats when model excludes the constant

$$BH_i^e + 99 = \beta_1 VWM_i^e + \beta_2 SMB_i + \beta_3 HML_i + \beta_4 MOM_i + \epsilon_i$$

```
In [34]: ls_p99 = smf.ols("I(BHe + 99) ~ VWMe + SMB + HML + MOM - 1", data).fit(cov_type="HC0")
summary(ls_lhs, [0, 1])
```

Dep. Variable:	I(BHe - VWMe - HML)	R-squared:	0.382			
Model:	OLS	Adj. R-squared:	0.378			
	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.0859	0.043	-1.991	0.046	-0.170	-0.001
VWMe	0.0798	0.012	6.910	0.000	0.057	0.102
SMB	0.0019	0.017	0.110	0.912	-0.032	0.036
HML	-0.2357	0.021	-11.219	0.000	-0.277	-0.195
MOM	-0.0354	0.013	-2.631	0.009	-0.062	-0.009

Estimating the residual variance

Small-sample corrected estimator

- Variance of shock estimated using model residuals

$$s^2 = \frac{1}{n - k} \sum_{i=1}^n \hat{\epsilon}_i^2 = \frac{\hat{\epsilon}' \hat{\epsilon}}{n - k}$$

```
In [36]: s2 = eps.T @ eps / (n - k)
pretty(f"{{s2:.3f}}")
```

1.132

Estimating the residual variance

Large-sample estimator

- Asymptotic results usually use the large-sample version of the variance estimator

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n \hat{\epsilon}_i^2 = \frac{\hat{\epsilon}' \hat{\epsilon}}{n}$$

```
In [37]: sigma2 = eps.T @ eps / n  
pretty(f"{{sigma2:.3f}}")
```

1.124

Scores and the first-order condition of OLS

- The FOC of a regression is

$$\mathbf{X}'\hat{\epsilon} = \sum_{i=1}^n \mathbf{x}_i' \hat{\epsilon}_i = \mathbf{0}$$

- Estimated residuals are *always* orthogonal with included regressors
- Later we will see these can be used to test models if $\approx \mathbf{0}$

```
In [39]: scores
```

```
Out[39]:
```

Scores	
Intercept	7.056578e-13
VWMe	1.274314e-11
SMB	5.496090e-13
HML	3.228973e-12
MOM	-2.858463e-12

Analysis of Cross-Sectional Data

Properties of OLS Estimators

- Invariance to Affine Transformations
- Asymptotic Distribution
- Feasible Central Limit Theorems
- Bootstrap Estimation of the Covariance

Variable Transformations

Rescaling by a constant

$$\frac{Y_i}{100} = \beta_1 + \beta_2 \frac{X_{i,2}}{100} + \dots + \beta_k \frac{X_{i,k}}{100} + \epsilon_i$$

```
In [40]: model = "BHe ~ 1 + VWMe + SMB + HML + MOM"
rescaled_ls = smf.ols(model, data / 100.0).fit(cov_type="HC0")
show_params(rescaled_ls, ls, columns=["Rescaled", "Orig"])
```

Out[40]:

	Rescaled	Orig
Intercept	-0.000859	-0.085899
VWMe	1.079785	1.079785
SMB	0.001894	0.001894
HML	0.764300	0.764300
MOM	-0.035397	-0.035397

Variable Transformations

Rescaling single variables

$$Y_i = \beta_1 + \beta_2 (2VWM_i^e) + \beta_3 SMB_i + \beta_4 \frac{HML_i}{2} + \beta_4 MOM_i + \epsilon_i$$

```
In [41]: model = "BHe ~ 1 + I(2 * VWMe) + SMB + I(1/2 * HML) + MOM"
ls_p10 = smf.ols(model, data).fit(cov_type="HC0")
show_params(ls_p10, columns=["Plus 10"])
```

Out[41]:

Plus 10	
Intercept	-0.085899
I(2 * VWMe)	0.539893
SMB	0.001894
I(1 / 2 * HML)	1.528600
MOM	-0.035397

Variable Transformations

Affine Transformations

$$(3BH^e + 7) = \beta_1 + \beta_2 (2VWM_i^e - 9) + \beta_3 \frac{SMB_i}{2} + \beta_4 HML_i + \beta_4 MOM_i + \epsilon_i$$

```
In [42]: model = "I(3 * BHe + 7) ~ 1 + I(2 * VWMe - 9) + I(1/2 *SMB) + HML + MOM"
ls_affine = smf.ols(model, data).fit(cov_type="HC0")
show_params(ls_affine, columns=["Affine"])
```

Out[42] :

Affine	
Intercept	21.319408
I(2 * VWMe - 9)	1.619678
I(1 / 2 * SMB)	0.011361
HML	2.292901
MOM	-0.106192

Characterizing Parameter Estimation Error

- Central Limit Theorem

$$\sqrt{n} \left(\hat{\beta}_n - \beta \right) \xrightarrow{d} N \left(\mathbf{0}, \Sigma_{XX}^{-1} \mathbf{S} \Sigma_{XX}^{-1} \right)$$

- Covariance components $\Sigma_{XX} = E \left[\mathbf{x}'_i \mathbf{x}_i \right]$ and $\mathbf{S} = \mathbf{p} -$

$$\lim_{n \rightarrow \infty} \text{Var} \left[\sqrt{n} \frac{1}{n} \sum_{i=1}^n \mathbf{x}'_i \epsilon_i \right]$$

- In practice

$$\hat{\beta}_n \approx N \left(\beta, \frac{\hat{\Sigma}_{XX}^{-1} \hat{\mathbf{S}} \hat{\Sigma}_{XX}^{-1}}{n} \right)$$

Characterizing Parameter Estimation Error

Parameter Covariance Matrix

In [44]: `ls.cov_params()`

Out[44]:

	Intercept	VWMe	SMB	HML	MOM
Intercept	0.001860	-0.000171	0.000079	-0.000157	-0.000154
VWMe	-0.000171	0.000133	-0.000060	0.000039	0.000019
SMB	0.000079	-0.000060	0.000297	0.000042	0.000019
HML	-0.000157	0.000039	0.000042	0.000441	0.000122
MOM	-0.000154	0.000019	0.000019	0.000122	0.000181

Characterizing Parameter Estimation Error

Estimating the Covariance

$$\hat{\Sigma}_{XX} = \frac{1}{n} \mathbf{X}' \mathbf{X} \text{ and } \hat{\mathbf{S}} = \sum_{i=1}^n \hat{\epsilon}_i^2 \mathbf{x}_i' \mathbf{x}_i$$

```
In [45]: xe = x * eps
S = xe.T @ xe / n
Sigma_inv = np.linalg.inv(x.T @ x / n)
cov = 1 / n * (Sigma_inv @ S @ Sigma_inv)
cov.index = cov.columns = x.columns
cov
```

Out[45] :

	Intercept	VWMe	SMB	HML	MOM
Intercept	0.001860	-0.000171	0.000079	-0.000157	-0.000154
VWMe	-0.000171	0.000133	-0.000060	0.000039	0.000019
SMB	0.000079	-0.000060	0.000297	0.000042	0.000019
HML	-0.000157	0.000039	0.000042	0.000441	0.000122
MOM	-0.000154	0.000019	0.000019	0.000122	0.000181

Characterizing Parameter Estimation Error

Standard Errors

- Root of diagonal elements of VCV

```
In [46]: pretty(ls.bse)
```

```
Out[46]:
```

Intercept	0.043134
VWMe	0.011547
SMB	0.017224
HML	0.021009
MOM	0.013455

Bootstrapping the Covariance

- Simulate from data to estimate covariance
- Randomly sample n observation with replacement (y_i, \mathbf{x}_i)
- Estimate $\hat{\beta}_b$ from random sample
- Repeat B times
- Compute covariance from bootstrapped $\hat{\beta}_b$

```
In [47]: betas = []
g = np.random.default_rng(2020)
lhs = ls.model.data.orig_endog
for i in range(1000):
    idx = g.integers(n, size=n)
    xb = x.iloc[idx]
    yb = lhs.iloc[idx]
    beta = sm.OLS(yb, xb).fit().params
    betas.append(beta)
betas = pd.DataFrame(betas, columns=x.columns, index=np.arange(1, len(betas) + 1))
betas.index.name = "b"
```

Boootstrap β estimates

```
In [48]: betas.head()
```

```
Out[48]:
```

	Intercept	VWMe	SMB	HML	MOM
b					
1	-0.032329	1.073407	-0.002263	0.747751	-0.059899
2	-0.098095	1.082409	0.047483	0.750626	-0.023940
3	-0.037763	1.085002	0.036935	0.776792	-0.025469
4	-0.019000	1.083847	-0.002470	0.689036	-0.053007
5	-0.052940	1.067186	0.032972	0.783204	-0.034621

Comparing the Bootstrap and the Traditional Estimator

In [49]: `betas.cov()`

Out[49]:

	Intercept	VWMe	SMB	HML	MOM
Intercept	0.001868	-0.000182	0.000124	-0.000100	-0.000134
VWMe	-0.000182	0.000135	-0.000062	0.000035	0.000015
SMB	0.000124	-0.000062	0.000313	0.000033	0.000018
HML	-0.000100	0.000035	0.000033	0.000405	0.000105
MOM	-0.000134	0.000015	0.000018	0.000105	0.000182

In [50]: `ls.cov_params()`

Out[50]:

	Intercept	VWMe	SMB	HML	MOM
Intercept	0.001860	-0.000171	0.000079	-0.000157	-0.000154
VWMe	-0.000171	0.000133	-0.000060	0.000039	0.000019
SMB	0.000079	-0.000060	0.000297	0.000042	0.000019
HML	-0.000157	0.000039	0.000042	0.000441	0.000122
MOM	-0.000154	0.000019	0.000019	0.000122	0.000181

Analysis of Cross-Sectional Data

Wald and t -tests

- Linear Equality Hypotheses
- Testing a Single Restriction with a t -tests
- The t -statistic
- Multiple Restrictions and the Wald tests
- The F -stats

Hypothesis Testing

- Null in a Linear Equality Test

$$H_0 : \mathbf{R}\beta = r$$

- Three classes of tests

- Wald and t -test
- Lagrange Multiplier
- Likelihood Ratio

Hypothesis Testing

t-tests

- Asymptotically normally distributed
- Test a single restriction
- Values outside of $\pm 1.96 \approx \pm 2$ lead to rejection using a 5% size
- Can be used to test 1-sided hypotheses

Hypothesis Testing

t-test Example

Testing the additional total effect is 0

$$H_0 : SMB + HML + MOM = 0$$
$$R = [0, 0, 1, 1, 1], r = 0$$

```
In [51]: summary(ls)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.0859	0.043	-1.991	0.046	-0.170	-0.001
VWMe	1.0798	0.012	93.514	0.000	1.057	1.102
SMB	0.0019	0.017	0.110	0.912	-0.032	0.036
HML	0.7643	0.021	36.380	0.000	0.723	0.805
MOM	-0.0354	0.013	-2.631	0.009	-0.062	-0.009

Hypothesis Testing

t-test Example

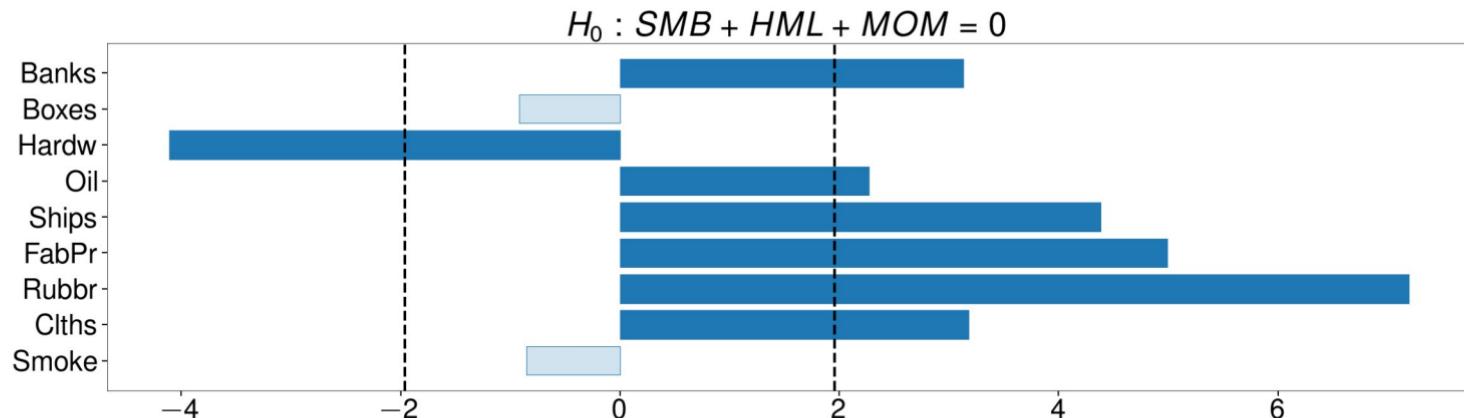
```
In [52]: R = np.array([[0, 0, 1, 1, 1]])
c = ls.cov_params()
h0_vcv = np.squeeze(R @ c @ R.T)
t = (R @ ls.params) / np.sqrt(h0_vcv)
pretty(f"The t-test statistic is {t[0]:0.2f}")
```

The t-test statistic is 20.39

Hypothesis Testing

t-test Example on Industry Portfolios

```
In [55]: test_plot(t_tests, title="H0: SMB + HML + MOM = 0", two_sided=True)
```



Hypothesis Testing

t-stats

- *t*-stat is special case for $H_0 : \beta_j = 0$
- Most commonly reported test statistic
- Asymptotic normal
- 5% critical values $\pm 1.96 \approx \pm 2$

```
In [56]: pretty(ls.tvalues)
```

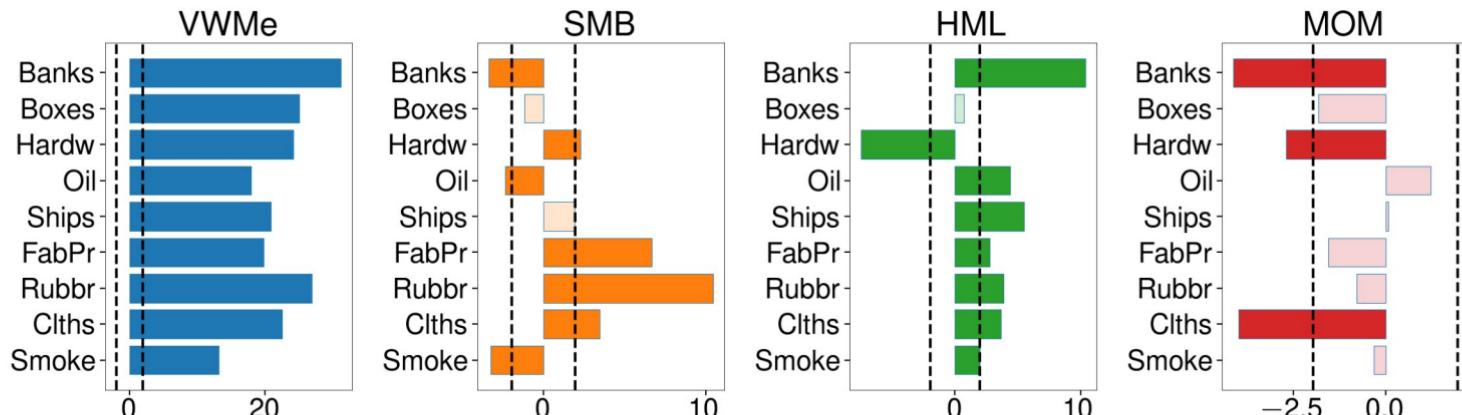
```
Out[56]:
```

Intercept	-1.991463
VWMe	93.513503
SMB	0.109934
HML	36.380381
MOM	-2.630803

Hypothesis Testing

Significance in Industry Portfolios

```
In [58]: multi_test_plot(t_stats)
```



Hypothesis Testing

Wald Tests

- Test multiple hypothesis
- Exploit properties of multivariate normals
- χ_m^2 distributed in large samples
- Test statistic is

$$W = n(\mathbf{R}\hat{\boldsymbol{\beta}} - \mathbf{r})' [\mathbf{R}\hat{\boldsymbol{\Sigma}}_{XX}^{-1}\hat{\mathbf{S}}\hat{\boldsymbol{\Sigma}}_{XX}^{-1}\mathbf{R}']^{-1} (\mathbf{R}\hat{\boldsymbol{\beta}} - \mathbf{r})$$

Hypothesis Testing

Wald Tests

Testing the CAPM

- Multiple β all zero: $H_0 : SMB$

$$= HML$$

$$= MOM = 0$$

$$R = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}, r = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Hypothesis Testing

Testing the CAPM

```
In [59]: R, r = np.zeros((3, 5)), np.zeros(3)
R[0, 2] = R[1, 3] = R[2, 4] = 1
h0_vcv = R @ c @ R.T
h0_vcv.columns = h0_vcv.index = [f"Restr {i}" for i in range(1, 4)]
h0_vcv
```

Out[59]:

	Restr 1	Restr 2	Restr 3
Restr 1	0.000297	0.000042	0.000019
Restr 2	0.000042	0.000441	0.000122
Restr 3	0.000019	0.000122	0.000181

```
In [60]: numerator = R @ ls.params - r
wald = numerator @ np.linalg.inv(h0_vcv) @ numerator.T
pretty(f"W={wald:.1f}")
```

W=1749.6

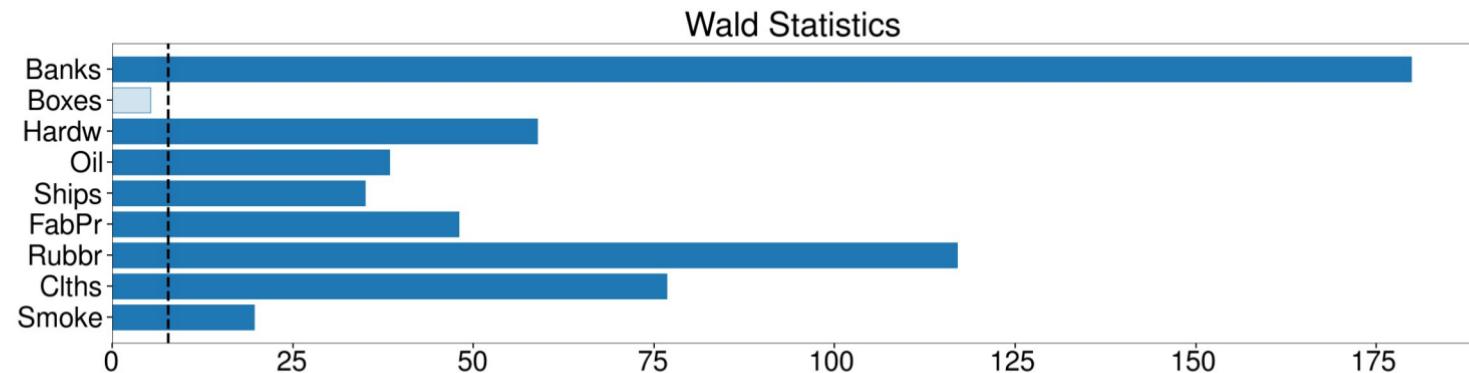
Wald Tests

Industry Portfolios

In [62]:

```
dof = 3
pretty(f"The crit. value is {stats.chi2(dof).ppf(0.95):0.2f} from a \chi^2_{dof}")
test_plot(walds, cv=stats.chi2(dof).ppf(0.95), title="Wald Statistics")
```

The crit. value is 7.81 from a χ^2_3



Hypothesis Testing

The F-stat

- Special case of Wald for

$$H_0 : \beta_2 = \beta_3 = \dots = \beta_k = 0$$

- Never test constant

- Restrictions in 4-factor model

$$R = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}, r = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

- If **no constant** then test all β

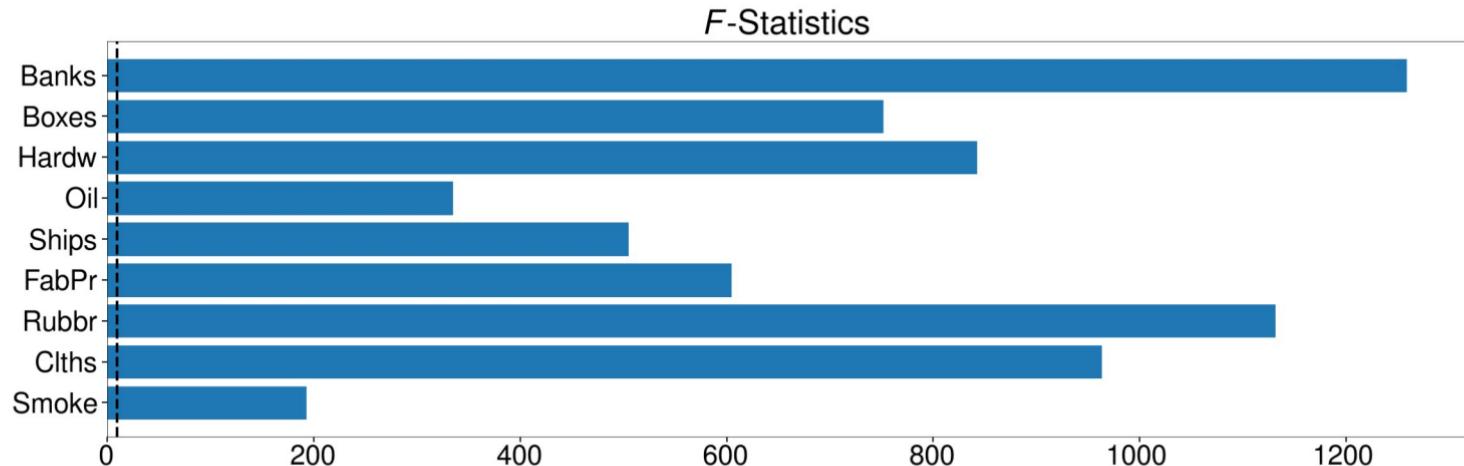
$$H_0 : \beta_1 = \beta_2 = \dots = \beta_k = 0$$

Industry F-statistics

In [64]:

```
dof = 4
cv = stats.chi2(dof).ppf(0.95)
pretty(f"The critical value is {cv:0.2f} from a \chi^2_{dof}")
test_plot(f_stats, cv=cv, title="$F$-Statistics")
```

The critical value is 9.49 from a χ_4^2



Analysis of Cross-Sectional Data

LM and LR Tests

- Imposing a LER on a Linear Regression
- LM Tests
- LR Tests
- Comparing Wald, LM and LR tests

Hypothesis Testing

Imposing the null on the model

$$H_0 : \beta_{SMB} + \beta_{HML} + \beta_{MOM} = 0 \Rightarrow \beta_{SMB} = -\beta_{HML} - \beta_{MOM}$$

Initial model

$$Ships = \beta_1 + \beta_2 VWM^e + \beta_3 SMB + \beta_4 HML + \beta_5 MOM + \epsilon_i$$

becomes

$$Ships = \beta_1 + \beta_2 VWM^e + (-\beta_4 - \beta_5) SMB + \beta_4 HML + \beta_5 MOM + \epsilon_i$$

and then finally

$$Ships = \beta_1 + \beta_2 VWM^e + \beta_4 (HML - SMB) + \beta_5 (MOM - SMB) + \epsilon_i$$

Hypothesis Testing

Estimating the Restricted Model

```
In [65]: model = "Ships ~ 1 + VWMe + I(HML-SMB) + I(MOM-SMB)"  
imposed = smf.ols(model, data).fit(cov_type="HC0")  
summary(imposed)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	0.0676	0.202	0.335	0.738	-0.328	0.463
VWMe	1.1401	0.058	19.805	0.000	1.027	1.253
I(HML - SMB)	0.1897	0.061	3.133	0.002	0.071	0.308
I(MOM - SMB)	-0.1304	0.061	-2.148	0.032	-0.249	-0.011

Hypothesis Testing

Lagrange Multiplier (LM) tests

- Uses property that scores should be 0 when model is correct
- Define scores

$$s_i = \mathbf{x}_i \tilde{\epsilon}_i$$

- Estimate score covariance using

$$\hat{S} = \frac{1}{n} \sum_{i=1}^n s_i' s_i$$

- LM test statistic is defined

$$LM = n \bar{s} \hat{S} \bar{s}' \xrightarrow{d} \chi_m^2$$

Hypothesis Testing

The Scores Using the Restricted Residuals

```
In [66]: imposed_eps = imposed.resid.to_numpy()
scores = x * imposed_eps[:, None]
mean_scores = scores.mean()
pretty(mean_scores)
```

Out[66]:

Intercept	-1.437140e-16
VWMe	1.680030e-14
SMB	1.614948e+00
HML	1.614948e+00
MOM	1.614948e+00

```
In [67]: S = scores.T @ scores / n
LM = n * mean_scores @ np.linalg.inv(S) @ mean_scores
pretty(f"The LM test statistic is {LM:.2f}")
```

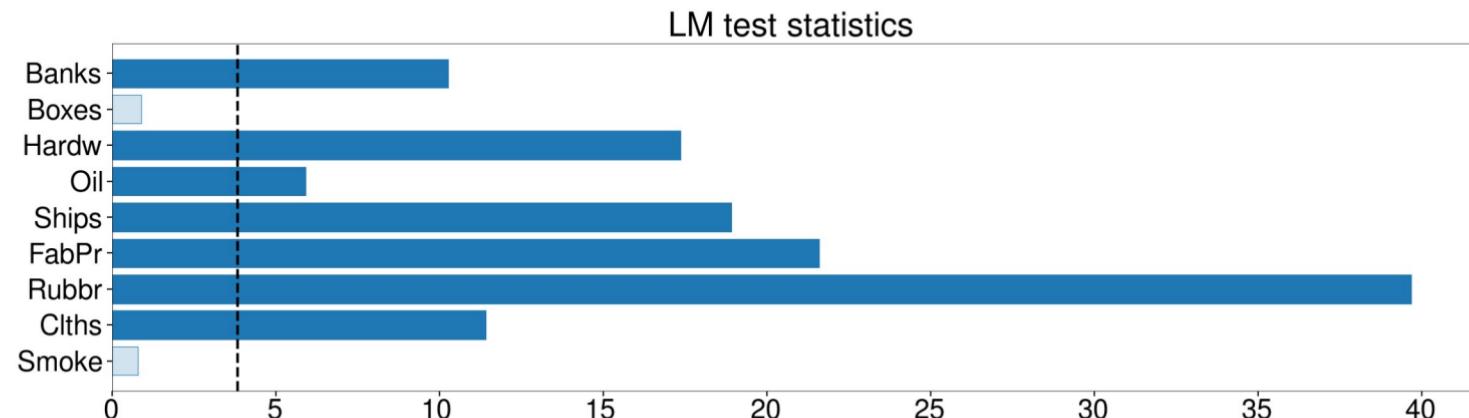
The LM test statistic is 18.91

Hypothesis Testing

LM Tests on Industry Portfolios

```
In [69]: cv = stats.chi2(1).ppf(0.95)
pretty(f"The critical value is {cv:.2f} from a \chi^2_1")
test_plot(lms, cv=cv, title="LM test statistics")
```

The critical value is 3.84 from a χ_1^2



Hypothesis Testing

Likelihood Ratio (LR) tests

- Nearly identical to LM, only using unrestricted model to estimate score covariance

$$\hat{s}_i = \mathbf{x}_i \hat{\epsilon}_i$$

- Covariance uses $\hat{\epsilon}_i$ instead of $\tilde{\epsilon}_i$

$$\hat{S} = \frac{1}{n} \sum_{i=1}^n \hat{s}_i' \hat{s}_i$$

- LR test statistic is defined

$$LR = n\bar{s}\hat{S}\bar{s}' \xrightarrow{d} \chi_m^2$$

Hypothesis Testing

Likelihood Ratio (LR) tests

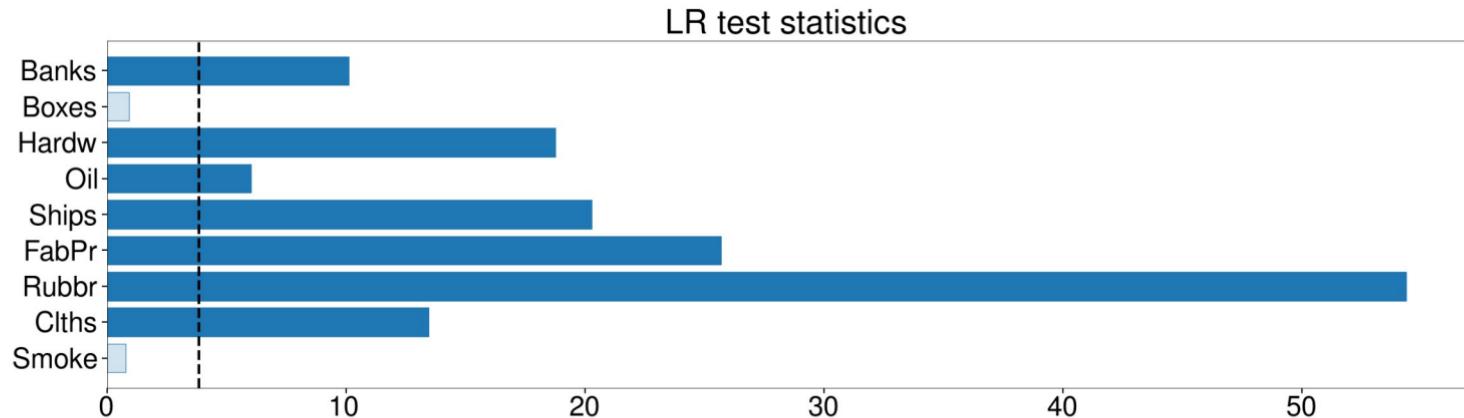
```
In [70]: unres = smf.ols("Ships ~ 1 + VWMe + SMB + HML + MOM", data).fit()
eps = unres.resid.to_numpy()
s_hat = x * eps[:, None]
S = s_hat.T @ s_hat / n
LR = n * mean_scores @ np.linalg.inv(S) @ mean_scores
pretty(f"The LR test statistic is {LR:.2f}")
```

The LR test statistic is 4.19

Hypothesis Testing

LM Tests on Industry Portfolios

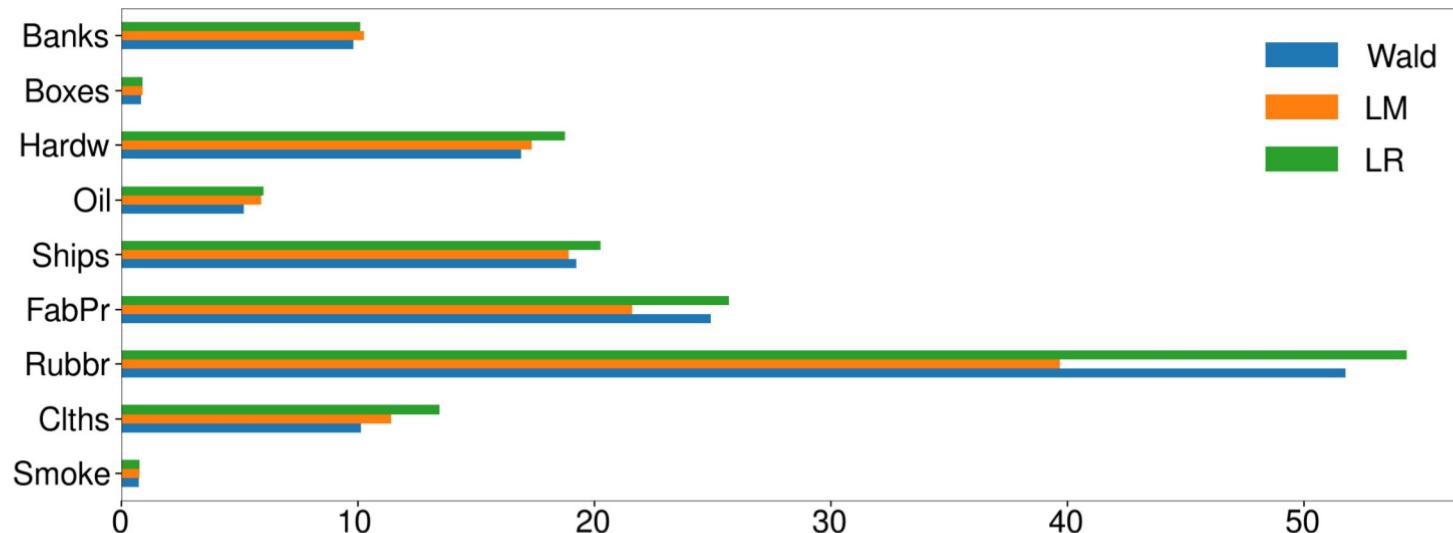
```
In [72]: test_plot(lrs, cv=cv, title="LR test statistics")
```



Hypothesis Testing

Comparing the Three Classes of Test

```
In [74]: plot_three_stats()
```



Analysis of Cross-Sectional Data

Heteroskedasticity

- Testing for Heteroskedasticity
- Covariance Estimation for Homoskedastic Data
- Bootstrap Covariance Estimation for Homoskedastic Data

Testing for Heteroskedasticity

White's test

- Key insight of White: Heteroskedasticity robust estimator only needed when

$$E[\epsilon_i^2 X_{i,o} X_{i,p}] \neq \sigma^2 E[X_{i,o} X_{i,p}]$$

- Use a regression to test if covariances are all 0

$$\hat{\epsilon}_i^2 = \mathbf{z}_i \boldsymbol{\gamma} + \eta_i$$

- \mathbf{z} contains all distinct cross-products of X

- Test statistic is nR^2 from auxiliary model

- $\chi^2_{k(k+1)/2-1}$ distribution when initial model includes a constant

Testing for Heteroskedasticity

White's test Implementation

```
In [75]: data["eps2"] = ls.resid ** 2

crosses = ""
for x1 in ("VWMe", "SMB", "HML", "MOM"):
    for x2 in ("VWMe", "SMB", "HML", "MOM"):
        crosses += f"+ I({x1} * {x2})"
formula = "eps2 ~ 1 + VWMe + SMB + HML + MOM" + crosses
pretty(formula)
```

$\text{eps2} \sim 1 + \text{VWMe} + \text{SMB} + \text{HML} + \text{MOM} + I(\text{VWMe} * \text{VWMe}) + I(\text{VWMe} * \text{SMB}) + I(\text{VWMe} * \text{HML}) + I(\text{VWMe} * \text{MOM}) + I(\text{SMB} * \text{VWMe}) + I(\text{SMB} * \text{SMB}) + I(\text{SMB} * \text{HML}) + I(\text{SMB} * \text{MOM}) + I(\text{HML} * \text{VWMe}) + I(\text{HML} * \text{SMB}) + I(\text{HML} * \text{HML}) + I(\text{HML} * \text{MOM}) + I(\text{MOM} * \text{VWMe}) + I(\text{MOM} * \text{SMB}) + I(\text{MOM} * \text{HML}) + I(\text{MOM} * \text{MOM})$

Testing for Heteroskedasticity

White's test Implementation

```
In [76]: white = smf.ols(formula, data).fit()  
summary(white, [0])
```

Dep. Variable:	eps2	R-squared:	0.109
Model:	OLS	Adj. R-squared:	0.090

Testing for Heteroskedasticity

White's test Regression

- Only showing significant coefficients

```
In [78]: summary_white()
```

	Intercept	VWMe	SMB	HML * HML	HML * MOM	MOM * HML	MOM * MOM
Parameter	8.134412e-01	-0.044663	0.063097	0.023196	0.010208	0.010208	0.003795
t-test Stat.	8.358941e+00	-2.275155	2.089132	4.640325	4.133473	4.133473	2.079526
p-value	3.645328e-16	0.023211	0.037072	0.000004	0.000040	0.000040	0.037948

Testing for Heteroskedasticity

White's test statistic

```
In [79]: white_stat = n * white.rsquared  
pretty(f"White's stat: {white_stat:0.2f}")
```

White's stat: 74.76

```
In [80]: pretty(f"Number of restrictions: {5 * (5 + 1) // 2 + 1}")
```

Number of restrictions: 16

```
In [81]: pvalue = 1.0 - stats.chi2(5 * (5 + 1) / 2 + 1).cdf(white_stat)  
pretty(f"The p-value is {pvalue:0.3f}")
```

The p-value is 0.000

Characterizing Parameter Estimation Error

Homoskedastic Data

- Central Limit Theorem when residuals are homoskedastic

$$\sqrt{n} \left(\hat{\boldsymbol{\beta}}_n - \boldsymbol{\beta} \right) \xrightarrow{d} N \left(\mathbf{0}, \sigma^2 \boldsymbol{\Sigma}_{XX}^{-1} \right)$$

- Covariance components $\boldsymbol{\Sigma}_{XX} = E \left[\mathbf{x}_i' \mathbf{x}_i \right]$ and $\sigma^2 = E[\epsilon_i^2]$.
- In practice

$$\hat{\boldsymbol{\beta}}_n \approx N \left(\boldsymbol{\beta}, \frac{\hat{\sigma}^2 \hat{\boldsymbol{\Sigma}}_{XX}^{-1}}{n} \right)$$

Characterizing Parameter Estimation Error

The Homoskedastic Parameter Covariance Estimator

```
In [83]: ls_homo = smf.ols("Smoke ~ 1 + VWMe + SMB + HML + MOM", data).fit()
summary(ls_homo)
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.5512	0.209	2.642	0.008	0.142	0.961
VWMe	0.7334	0.049	14.917	0.000	0.637	0.830
SMB	-0.2761	0.070	-3.948	0.000	-0.413	-0.139
HML	0.1671	0.075	2.243	0.025	0.021	0.313
MOM	-0.0204	0.050	-0.407	0.684	-0.119	0.078

Comparing Std Errors and t -stats

In [85]: `compare_tstats()`

Out[85]:

	Homoskedastic			Heteroskedastic	
	param	std err	t	std err	t
Intercept	0.551242	0.208643	2.642032	0.043134	-1.991463
VWMe	0.733354	0.049162	14.917088	0.011547	93.513503
SMB	-0.276084	0.069932	-3.947901	0.017224	0.109934
HML	0.167094	0.074507	2.242654	0.021009	36.380381
MOM	-0.020364	0.049996	-0.407308	0.013455	-2.630803

Characterizing Parameter Estimation Error

The Homoskedastic Parameter Covariance Estimator

```
In [86]: ls_homo.cov_params()
```

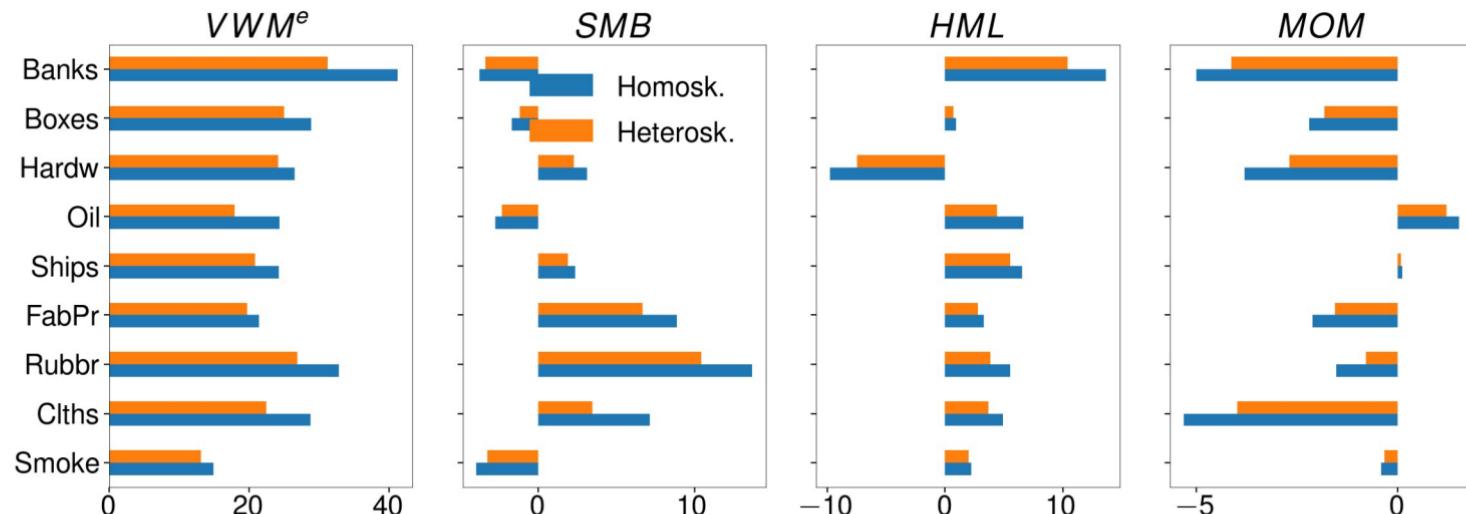
```
Out[86]:
```

	Intercept	VWMe	SMB	HML	MOM
Intercept	0.043532	-0.001698	-0.000538	-0.002558	-0.002149
VWMe	-0.001698	0.002417	-0.000916	0.000814	0.000477
SMB	-0.000538	-0.000916	0.004890	0.000589	0.000019
HML	-0.002558	0.000814	0.000589	0.005551	0.000882
MOM	-0.002149	0.000477	0.000019	0.000882	0.002500

Heteroskedasticity vs Homoskedasticity

Industry Portfolios

```
In [88]: plot_tvalues()
```



Bootstrap for Homoskedastic Data

- If data are homoskedastic can use improved bootstrap
- Independently sample $\hat{\epsilon}_i$ and \mathbf{x}_j and then build simulated $\tilde{Y}_m = \mathbf{x}_j \hat{\beta} + \hat{\epsilon}_i$
- Estimate model on bootstrapped data
- Repeat $b = 1, 2, \dots, B$ times and compute covariance of estimated $\hat{\beta}_b$

Bootstrap for Homoskedastic Data

In [89]:

```
eps = ls.resid.to_numpy()
betas = []
g = np.random.default_rng(2020)
x = ls.model.data.orig_exog
for i in range(1000):
    x_idx = g.integers(n, size=n)
    xb = x.iloc[x_idx]
    eps_idx = g.integers(n, size=n)
    y = xb @ ls.params + eps[eps_idx]
    beta = sm.OLS(y, xb).fit().params
    betas.append(beta)
betas = np.array(betas)

betas = pd.DataFrame(betas, columns=x.columns)
```

Bootstrap for Homoskedastic Data

In [90]: `betas.cov()`

Out[90]:

	Intercept	VWMe	SMB	HML	MOM
Intercept	0.001796	-0.000077	0.000007	-0.000104	-0.000095
VWMe	-0.000077	0.000098	-0.000030	0.000032	0.000013
SMB	0.000007	-0.000030	0.000198	0.000026	0.000011
HML	-0.000104	0.000032	0.000026	0.000229	0.000040
MOM	-0.000095	0.000013	0.000011	0.000040	0.000103

Analysis of Cross-Sectional Data

Model Selection

- General-to-Specific
- Specific-to-General
- Information Criteria
- Cross-Validation

Model Selection

General-to-Specific

- Start with full model
- Drop variables one-at-a-time when P-value > α
 - Typical sizes 1% or .1%

```
In [91]: res = smf.ols("Ships ~ 1 + VWMe + SMB + HML + MOM", data).fit(cov_type="HC0")
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1420	0.204	-0.697	0.486	-0.541	0.257
VWMe	1.1551	0.055	20.872	0.000	1.047	1.264
SMB	0.1604	0.084	1.898	0.058	-0.005	0.326
HML	0.4703	0.085	5.513	0.000	0.303	0.637
MOM	0.0055	0.074	0.074	0.941	-0.140	0.151

General-to-Specific Model Selection

```
In [92]: res = smf.ols("Ships ~ 1 + VWMe + SMB + HML", data).fit(cov_type="HC0")
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1372	0.197	-0.697	0.486	-0.523	0.248
VWMe	1.1541	0.056	20.716	0.000	1.045	1.263
SMB	0.1603	0.084	1.900	0.057	-0.005	0.326
HML	0.4683	0.086	5.474	0.000	0.301	0.636

General-to-Specific Model Selection

```
In [93]: res = smf.ols("Ships ~ 1 + VWMe + HML", data).fit(cov_type="HC0")
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1201	0.197	-0.610	0.542	-0.506	0.266
VWMe	1.1842	0.054	21.870	0.000	1.078	1.290
HML	0.4492	0.084	5.371	0.000	0.285	0.613

```
In [94]: res = smf.ols("Ships ~ VWMe + HML - 1", data).fit(cov_type="HC0")
summary(res)
```

	coef	std err	z	P> z	[0.025	0.975]
VWMe	1.1802	0.054	22.045	0.000	1.075	1.285
HML	0.4442	0.083	5.332	0.000	0.281	0.608

Model Selection

Specific-to-General

- Start with only a constant
- Add variables one-at-a-time and keep smallest P-value if $> \alpha$

```
In [95]: excl = ["VWMe", "SMB", "HML", "MOM"]
for reg in excl:
    res = smf.ols(f"Ships ~ 1 + {reg}", data).fit(cov_type="HC0")
    print(f"{reg}: {res.pvalues[reg]:0.3f}")
```

VWMe: 0.000
SMB: 0.000
HML: 0.793
MOM: 0.014

Specific-to-General

Removing insignificant variables

```
In [96]: excl.remove("VWMe")
for reg in excl:
    res = smf.ols(f"Ships ~ 1 + VWMe + {reg}", data).fit(cov_type="HC0")
    print(f"{reg}: {res.pvalues[reg]:0.3f}")
```

SMB: 0.263
HML: 0.000
MOM: 0.386

```
In [97]: excl.remove("HML")
for reg in excl:
    res = smf.ols(f"Ships ~ 1 + VWMe + HML + {reg}", data).fit(cov_type="HC0")
    print(f"{reg}: {res.pvalues[reg]:0.3f}")
```

SMB: 0.057
MOM: 0.947

Model Selection

Information Criteria

- Information criteria trade-off fit and cost for additional penalties
- Two most common: AIC and BIC
- Select model that produces the smallest IC from candidate models

```
In [98]: capm = smf.ols("Ships ~ 1 + VWMe", data).fit()
factor2 = smf.ols("Ships ~ 1 + VWMe + SMB", data).fit()
pretty(f"CAPM AIC: {capm.aic:0.1f}, BIC: {capm.bic:0.1f}")
pretty(f"2 Factor AIC: {factor2.aic:0.1f}, BIC: {factor2.bic:0.1f}")
```

CAPM AIC: 4225.9, BIC: 4234.9

2 Factor AIC: 4225.4, BIC: 4239.0

Model Selection

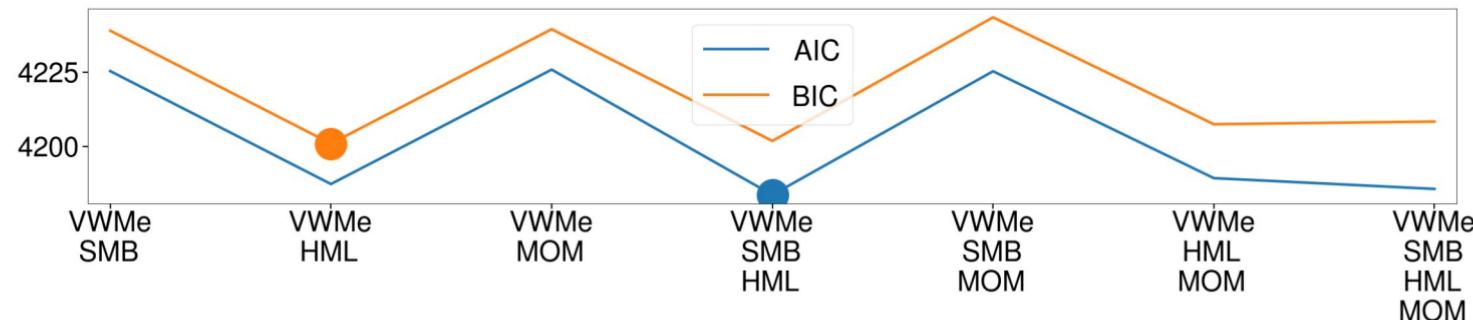
Global IC Search

```
In [100]: pretty(ics.idxmin())
```

```
Out[100]:
```

AIC	VWMe,SMB,HML
BIC	VWMe,HML

```
In [101]: plot_ics()
```



Model Selection

k -fold Cross-validation

- Focus on pseudo-out-of-sample prediction
- Split data into k equally sized random blocks
- Estimate parameters using $k - 1$ blocks
- Evaluate SSE using block not used in estimation
- Repeat k times in total computing the SSE once in each block
- Sum k SSE values into SSE_{xv}
- Choose model with the lowest out-of-sample SSE_{xv}

Model Selection

k-fold Cross-validation in Practice

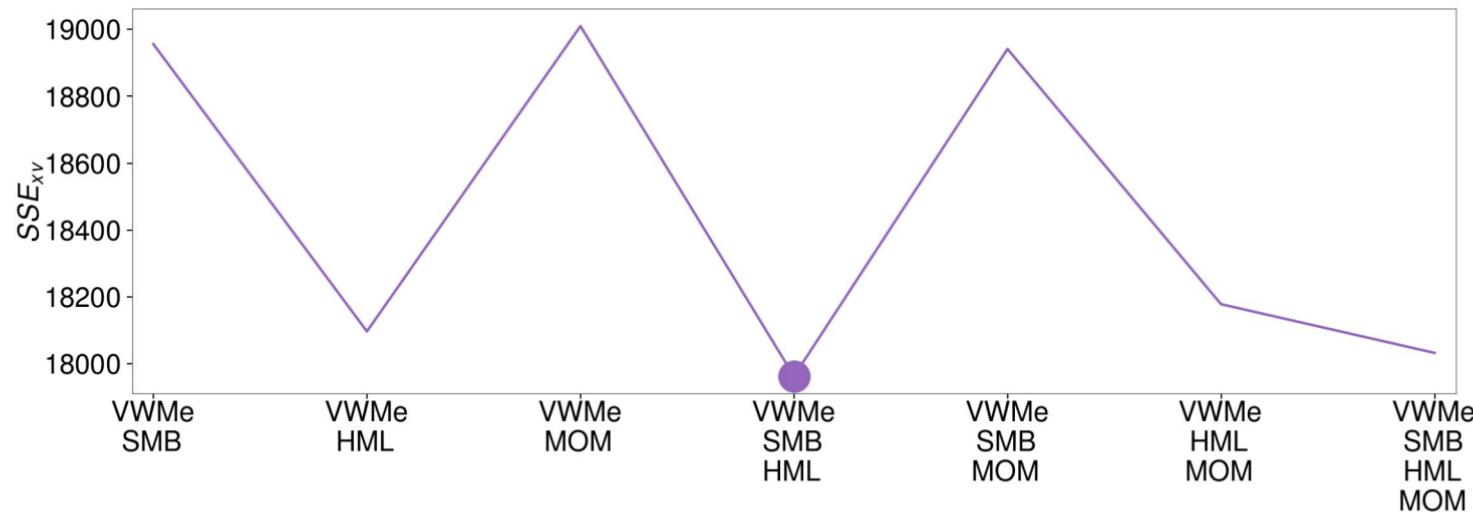
```
In [102]: mod = "Ships ~ 1 + VWMe + SMB + HML + MOM"
rg = np.random.default_rng(13221711120)
idx = rg.permutation(n)
fifth = n / 5
xv_errors = data.Ships.copy()
for i in range(5):
    reserve = idx[int(i * fifth) : int((i + 1) * fifth)]
    use = np.setdiff1d(idx, reserve)
    beta = smf.ols(mod, data.iloc[use]).fit().params
    xv_predictions = smf.ols(mod, data.iloc[reserve]).predict(beta)
    xv_errors.iloc[reserve] = data.Ships.iloc[reserve] - xv_predictions
sse_xv = (xv_errors ** 2).sum()
full_res = smf.ols(mod, data).fit()
pretty(f"XV SSE: {sse_xv:0.1f}, In-sample SSE: {n*full_res.mse_resid:0.1f}")
```

XV SSE: 18032.7, In-sample SSE: 17809.5

Model Selection

Cross-validation

```
In [104]: xv_plot()
```



Analysis of Cross-Sectional Data

Checking for Specification Errors

- Testing Structural Stability
- Testing for Neglected Nonlinearities
- Visual Diagnostics
- Trimming and Winsorization

Specification Testing

The Chow Test

- Chow test is a stability test
- Implemented using dummy interaction variables

$$I_{[t>\tau]}$$

- Extend model with copy of variables interacted

$$Y_t = \mathbf{x}_t \boldsymbol{\beta} + I_{[t>\tau]} \mathbf{x}_t \boldsymbol{\gamma} + \epsilon_t$$

- Test using a Wald test (or LM or LR) with a χ^2_k distribution

The Chow Test

```
In [105]: chow_mod = "Banks ~ 1 + VWMe + SMB + HML + MOM + IxVWMe + IxSMB + IxHML + IxMOM"
ind = data.index > pd.to_datetime("1987-10-1")
interact = data[["VWMe", "SMB", "HML", "MOM"]] * ind[:, None]
interact.columns = [f"Ix{col}" for col in interact]
both = pd.concat([data, interact], 1)
chow = smf.ols(chow_mod, both).fit(cov_type="HC0")
summary(chow, [1])
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.0588	0.126	-0.465	0.642	-0.307	0.189
VWMe	1.1137	0.057	19.469	0.000	1.002	1.226
SMB	-0.0774	0.072	-1.081	0.280	-0.218	0.063
HML	0.2557	0.086	2.990	0.003	0.088	0.423
MOM	-0.2023	0.059	-3.428	0.001	-0.318	-0.087
IxVWMe	0.0683	0.073	0.942	0.346	-0.074	0.211
IxSMB	-0.0842	0.089	-0.942	0.346	-0.259	0.091
IxHML	0.5119	0.103	4.947	0.000	0.309	0.715
IxMOM	0.0977	0.069	1.422	0.155	-0.037	0.232

The Chow Test

- Restrict the interaction terms to be 0

```
In [106]: R = np.c_[np.zeros((4, 5)), np.eye(4)]  
R
```

```
Out[106]: array([[0., 0., 0., 0., 0., 1., 0., 0., 0.],  
                  [0., 0., 0., 0., 0., 0., 1., 0., 0.],  
                  [0., 0., 0., 0., 0., 0., 0., 1., 0.],  
                  [0., 0., 0., 0., 0., 0., 0., 0., 1.]])
```

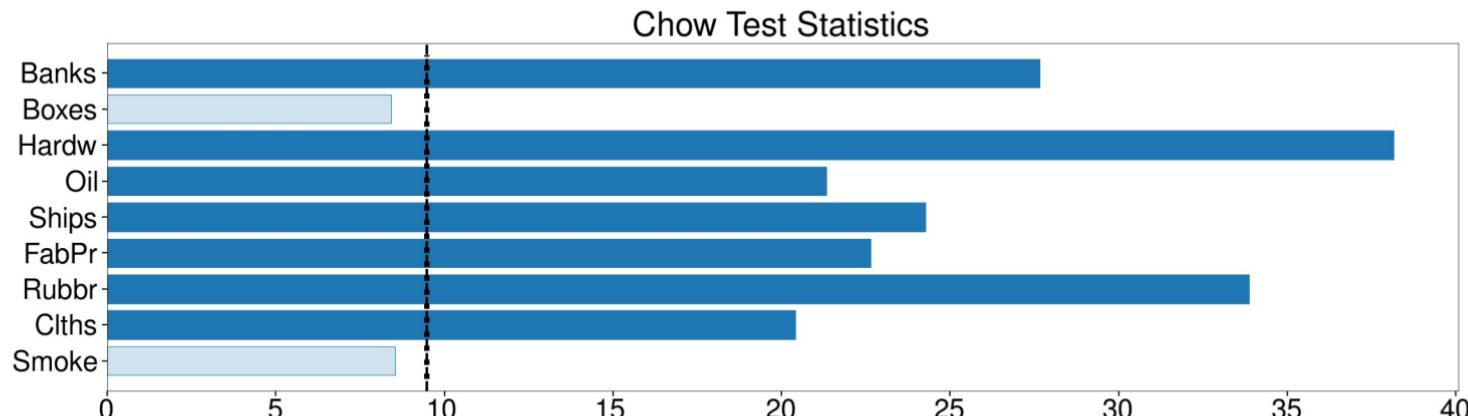
```
In [107]: chow_test = chow.wald_test(R)  
chow_stat = float(chow_test.statistic)  
chow_pvalue = chow_test.pvalue  
pretty(f"The Chow statistic is {chow_stat:.1f} and its p-value is {chow_pvalu  
e:.3f}")
```

The Chow statistic is 27.7 and its p-value is 0.000

Chow Test

Industry Portfolios

```
In [109]: cv = stats.chi2(4).ppf(0.95)
test_plot(chows, title="Chow Test Statistics", cv=cv)
_ = plt.plot([cv, cv], [-0.5, 8.5], "k:", linewidth=8)
```



Specification Testing

The RESET Test

- Test for general neglected nonlinearity
- Include powers of fitted value \hat{Y}_i^p in the model
- Requires initial regression to generate fitted value $Y_i = \mathbf{x}_i\boldsymbol{\beta} + \gamma_2\hat{Y}_i^2 + \gamma_3\hat{Y}_i^3 + \epsilon_i$

```
In [110]: first_stage = smf.ols("FabPr ~ 1 + VWMe + SMB + HML + MOM", data).fit()
data["FabPrHat"] = first_stage.predict()
reset_res = smf.ols("FabPr ~ 1 + VWMe + SMB + HML + MOM + I(FabPrHat**2)", dat
a).fit()
test_and_pval = pd.concat([reset_res.tvalues, reset_res.pvalues], 1)
test_and_pval.columns = ["t-stat", "p-value"]
```

The RESET Test Regression Results

```
In [111]: test_and_pval
```

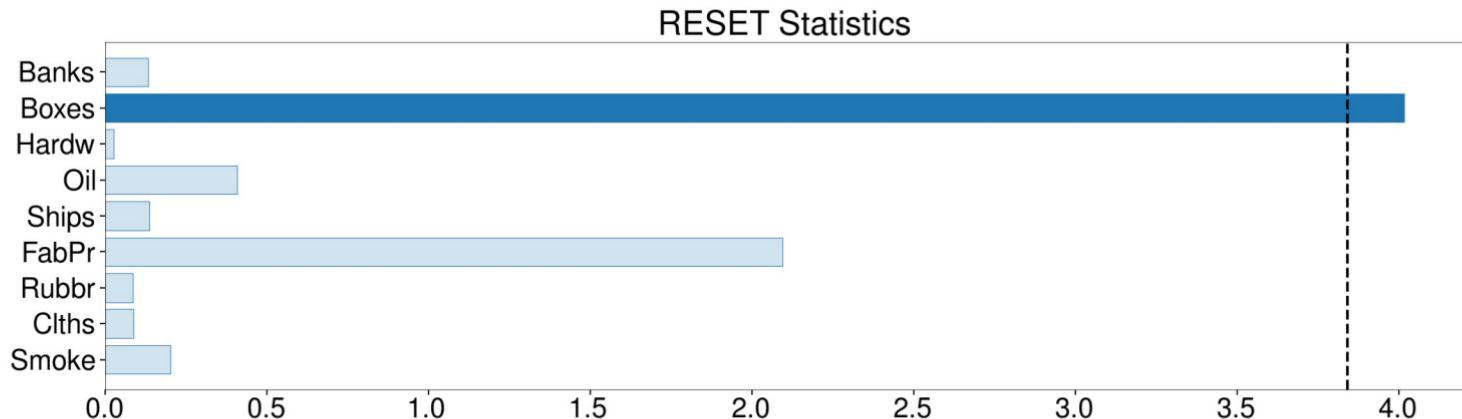
```
Out[111]:
```

	t-stat	p-value
Intercept	-0.406132	6.847736e-01
VWMe	20.341804	3.244811e-72
SMB	8.877870	6.004409e-18
HML	3.095003	2.048885e-03
MOM	-2.513712	1.217672e-02
I(FabPrHat ** 2)	-1.953090	5.121869e-02

The RESET Test

Industry Portfolios

```
In [113]: cv = stats.chi2(1).ppf(0.95)
test_plot(reset_stats, title="RESET Statistics", cv=cv)
```



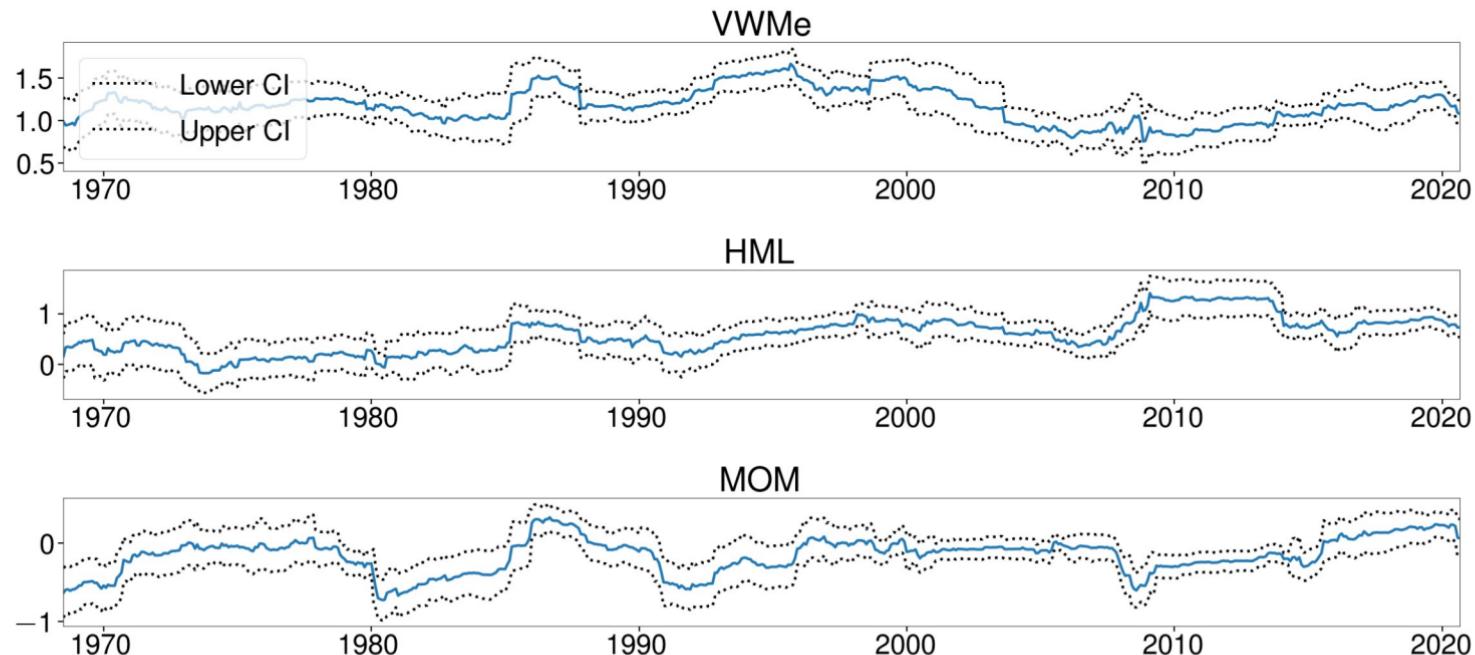
Specification Testing

Rolling Regression Plots

- Estimate many regressions using a fixed window length
- Visual diagnostics for parameter stability
- Confidence intervals are approximate

Rolling Regression Plots

```
In [115]: rolling_plot()
```



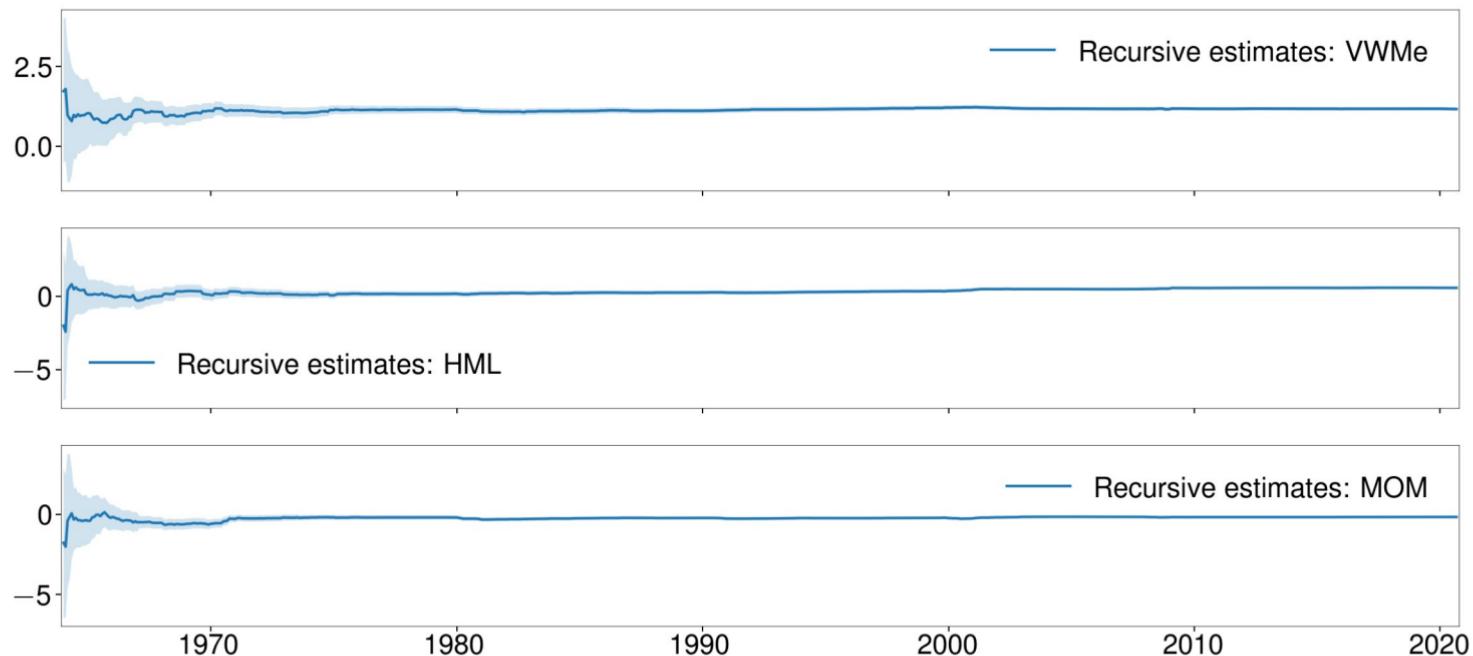
Specification Testing

Recursive Regression Plots

- Estimate many regressions using an expanding sample
- Visual diagnostics for parameter stability
- Confidence intervals are standard

Recursive Regression Plots

```
In [117]: recursive_plot()
```



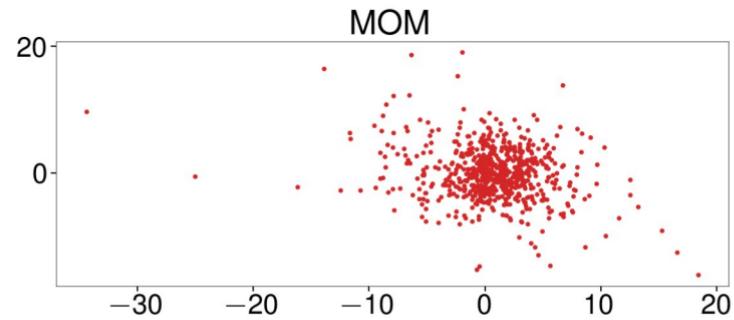
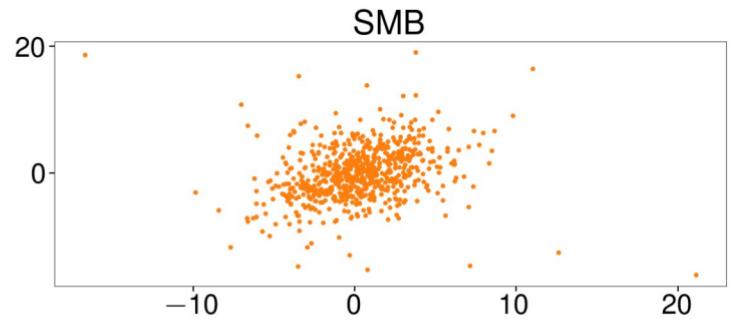
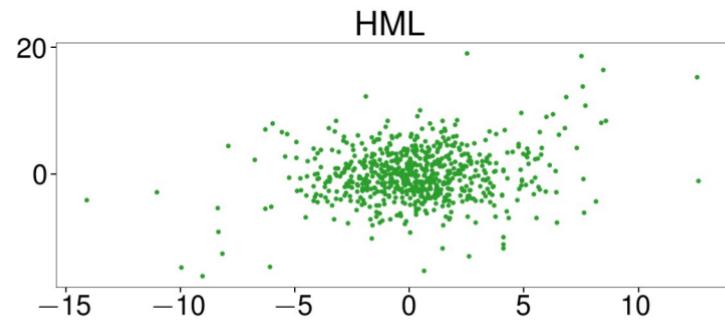
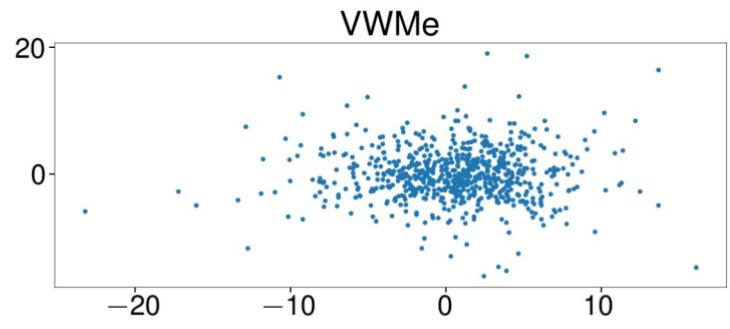
Specification Testing

Residual Plots: Residual vs X

- Residual plots are simple method to detect visible misspecification

Residual Plots: Residual vs X

```
In [119]: plot_residual()
```



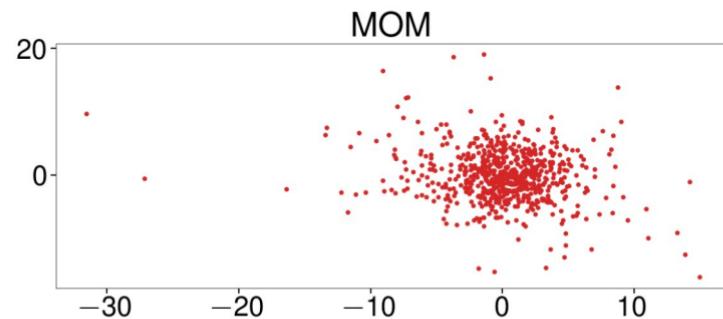
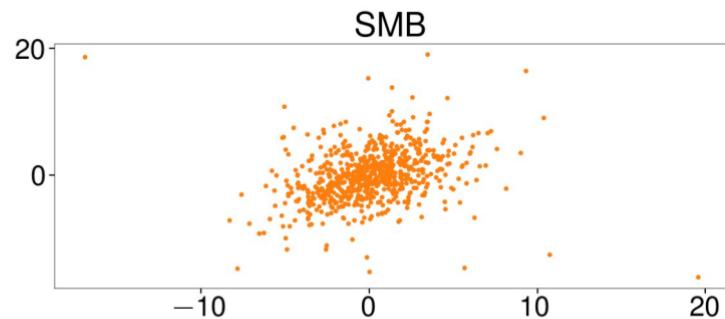
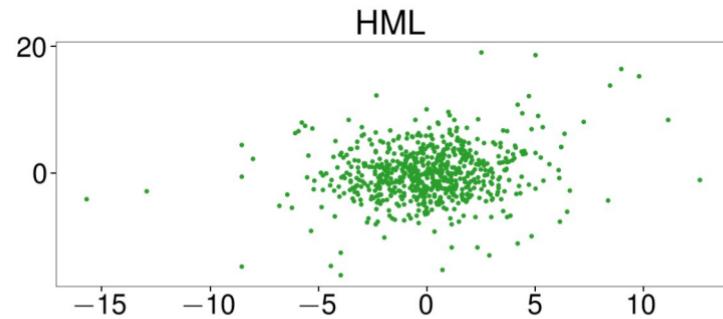
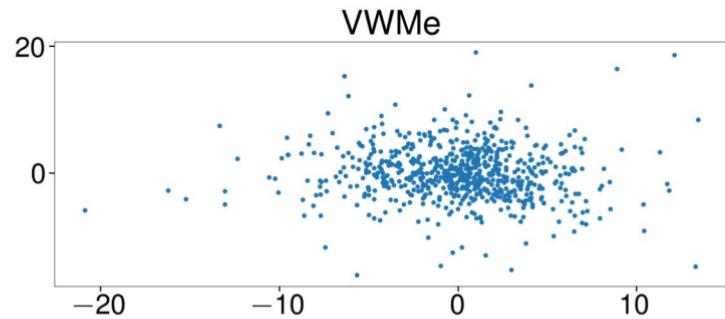
Specification Testing

Residual Plots: Residual vs $X_i | X_1, X_2, \dots, X_{i-1}, X_{i+1}, \dots, X_k$

- Sometimes useful to plot residuals against regressors *partialed* out
- Regress each regressor on the others included in the model
- Residual contains unique component of each regressor

Partial Residual Plot

```
In [121]: plot_partial_residual()
```



Trimming

- Drop observations with large ϵ_i
- Requires an initial estimator of β
 - "Good" subset of data
 - Typical value in random subset
 - Robust estimator or LAD
- Remove observations with $\hat{\epsilon}_i$ below quantile α and above $1 - \alpha$ for small α (1%, 2.5%, 5%)

```
In [122]: lad = smf.quantreg("Oil ~ 1 + VWMe + SMB + HML + MOM", data).fit(q=0.5)
bounds = lad.resid.quantile([0.025, 0.975])
pretty(bounds)
```

Out[122]:

0.025	-7.929285
0.975	8.650551

Trimming

```
In [123]: retain = (lad.resid > bounds.iloc[0]) & (lad.resid < bounds.iloc[1])
mod = "Oil ~ 1 + VWMe + SMB + HML + MOM"
trimmed = smf.ols(mod, data.loc[retain]).fit(cov_type="HC0")
summary(trimmed)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1302	0.141	-0.925	0.355	-0.406	0.146
VWMe	0.9070	0.038	23.716	0.000	0.832	0.982
SMB	-0.1617	0.055	-2.918	0.004	-0.270	-0.053
HML	0.3853	0.062	6.224	0.000	0.264	0.507
MOM	0.0567	0.039	1.448	0.148	-0.020	0.133

Compare to OLS Estimates

```
In [124]: full = smf.ols("Oil ~ 1 + VWMe + SMB + HML + MOM", data).fit(cov_type="HC0")
summary(full)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1240	0.172	-0.722	0.470	-0.461	0.213
VWMe	0.9537	0.053	17.951	0.000	0.850	1.058
SMB	-0.1525	0.065	-2.335	0.020	-0.281	-0.024
HML	0.3944	0.089	4.414	0.000	0.219	0.569
MOM	0.0609	0.050	1.218	0.223	-0.037	0.159

Windsorization

- Similar to trimming with one key difference
- Replace large ϵ_i with their quantile
- No data dropped

```
In [125]: data["Windsorized"] = data.Oil
predicted = lad.predict()
low = lad.resid < bounds.iloc[0]
data.loc[low, "Windsorized"] = predicted[low] + bounds.iloc[0]
high = lad.resid > bounds.iloc[1]
data.loc[high, "Windsorized"] = predicted[high] + bounds.iloc[1]
mod = "Windsorized ~ 1 + VWMe + SMB + HML + MOM"
windsorized = smf.ols(mod, data).fit(cov_type="HC0")
summary(windsorized)
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-0.1163	0.154	-0.755	0.451	-0.419	0.186
VWMe	0.9225	0.042	21.844	0.000	0.840	1.005
SMB	-0.1519	0.059	-2.568	0.010	-0.268	-0.036
HML	0.3700	0.069	5.389	0.000	0.235	0.505
MOM	0.0564	0.043	1.303	0.193	-0.028	0.141